

IRSIM 9.7 Reference

Table of Contents

[IRSIM Usage](#)
[Online IRSIM Tutorial](#)
[IRSIM Standard Command Set](#)
[Analyzer GUI Interface Commands](#)
[Statistics Package Commands](#)
[Connection-List Statistics Package Commands](#)
[Fault Simulator Package Commands](#)
[Power Estimation Package Commands](#)
[User Subcircuit Package](#)
[Random Number Generator Package Commands](#)
[The Analyzer GUI Interface](#)
[Using IRSIM With Tcl/Tk-Based Magic](#)
[Technology \(.prm\) Files for IRSIM Version 9.7](#)
[Documentation for IRSIM](#)
[Specifying X resources for IRSIM](#)
[Troubleshooting IRSIM](#)

IRSIM Usage:

```
irsim [-s] [-p power_net] [-g ground_net] prm_file sim_file ... [+hist_file] [-tcl_file ...] [-c tcl_file ...] [-  
@ cmd_file ...]
```

where the switches and arguments are as follows:

-s

If specified, 2 or more transistors of the same type connected in series, with no other connections to their common source/drain will be stacked into a compound transistor with multiple gates.

-p *power_net*

If specified, the net named *power_net* is assumed to be the power supply. Simulation can be done without specifying a power net, but any undeclared power net must be manually set to logic one in the command file, and the time to parse the netlist and start up will be increased.

-g *ground_net*

If specified, the net named *power_net* is assumed to be the ground return. Simulation can be done without specifying a ground net, but any undeclared ground net must be manually set to logic zero in the command file, and the time to parse the netlist and start up will be increased.

prm_file

The electrical parameters file that configure the devices to be simulated. It defines the capacitance of the various layers, transistor resistances, threshold voltages, and so forth. If *prm_file* does not specify an absolute path then IRSIM will search for the *prm_file* as follows (in order):

1. *./prm_file* (in the current directory).
2. $\${CAD_ROOT}/irsim/prm_file$
3. $\${CAD_ROOT}/irsim/prm_file$

The name *prm_file* may or may not contain the usual file extension ".prm".

sim_file ...

All file names *not* beginning with a '-' are assumed to be sim (netlist) files. These are files in ".sim" format, a flat description of a circuit's transistors and parasitic passive devices (R and C). The name *sim_file* may or may not contain the usual file extension ".sim". Each netlist file is read in turn and added to the network database. There is only a single name space for nodes, so references to node "A" in different network files all refer to the same node. While this feature

allows one to modularize a large circuit into several network files, care must be taken to ensure that no unwanted node merges happen due to an unfortunate clash in names.

`[-tcl_file ...]`

File names prefaced with a '-' are assumed to be command files, text files which contain command lines to be processed in the normal fashion. These files are processed line by line. When an end-of-file is encountered, processing continues with the next file. After all the command files have been processed, and if an **"exit"** command has not terminated the simulation run, IRSIM will accept further commands from the command line.

`[-c tcl_file]`

This is an alternative syntax for specifying a command file. It must occur on the command line *after* any simulation (.sim) files. The alternative syntax allows UNIX shell command-line completion of the name *cmd_file*, which is disabled when the filename is connected to the "-" in front of it.

`[-@ cmd_file]`

This is another alternative form that forces a command file to be read using syntax that is backwards-compatible to non-Tcl versions of IRSIM, in which the commands "set" and "flush" are used, instead of their altered counterparts in the Tcl interpreter environment, "setvector" and "histflush". This usage is equivalent to starting IRSIM without specifying a command file, then typing "@ *cmd_file*" at the IRSIM prompt.

`[+hist_file]`

The *hist_file* is the name of a file created with the **dump** command (see below). If it is present, IRSIM will initialize the network to the state saved in that file. This file is different from the ones created with the ">" command since it saves the state of every node for all times, including any pending events.

[Online IRSIM Tutorial](#)

IRSIM Standard Command Set

IRSIM commands and their syntax are as follows:

<
see **restorestate**

<<
see **restoreall**

>
see **savestate**

!
see **querygate**

?
see **querysource**

@ filename

Read and execute simulation commands from the file *filename*. This command works in a backwards-compatible mode that allows files in the original command-file syntax to be read. The obsolete commands "set" and "flush", which are incompatible with Tcl syntax, are interpreted to mean "setvector" and "histflush", respectively. To read command files using Tcl commands "set" and/or "flush", use the Tcl command **"source"** (see below).

activity *from* [*to*]

Report circuit activity in time interval *from* to *to*, if specified, or the current simulation time, if unspecified.

addnode *name*

Adds a new node named *name* to the IRSIM database. This node is unconnected, but it can be driven to specific states with the 'h', 'l', and 'u' commands, added to vectors, and displayed in the analyzer. It is useful for generating timed reference events in the analyzer, or to import simulation events from another program.

alias [*nodename* [*nodealias* ...]]

print node aliases (if there are any). With one argument, prints the aliases of node *nodename* (again, if there are any). With two or more arguments, nodes named *nodealias* will be connected in the database.

ana *node|vector*

shorthand for **analyzera** (see below)

analyzer [-b|-o|-d|-h] *node|vector*...

display node/vector(s) in the analyzer window. This command causes the analyzer window to be initialized and displayed if it was not already created; otherwise, the indicated node or vector is added to the display. If the signal name is a vector, then the option switches force that signal to be displayed in the indicated numeric base ("-b" for binary, "-d" for decimal, "-o" for octal, and "-h" for hexadecimal).

assert *node|vector* [*mask*] *val*

Assert that the indicated node or vector is equal to value *val* after bit-masking with *mask*.

assert *node*

Short form. Without an assertion, returns the value of node *node*. This can be used directly in conjunction with magic to query the simulation value of nodes without adding them to the watch or display lists. To return the value of a layout node, select paint and do (verbatim):

assert [getnode]

assertWhen *nodeT valT node val*

assert that the value of *node* is equal to *val* upon the condition that node *nodeT* switches to value *valT*

at [+]*time* { *procedure* }

(command introduced in IRSIM version 9.7.32)

Schedule the command or set of commands "*procedure*" to be executed at time *time* (in nanoseconds). In the Tcl/Tk version of IRSIM, *procedure* may be any valid set of Tcl commands or IRSIM commands. The procedure is executed exactly once, at the specified time. If the *time* value is preceded by "+", then the time is interpreted as a time relative to the current simulation time. In the Tcl version of IRSIM, the **at** command returns a unique ID for the event that can be used for the **cancel** and **get** options (see below). Also see the command **every**.

at get|cancel *id*

(command introduced in IRSIM version 9.7.32)

The **get** option returns the procedure associated with the scheduled event whose identifying (integer) tag is *id*. The **cancel** option cancels the scheduled event whose identifying tag is *id*.

back *time*

move simulation time back to specified *time*

c [*n*]

simulate for *n* clock cycles (default 1) or continue last simulation command prior to stopping. Use of this command requires that a clock has been defined using the "**clock**" command.

changes *from* [*to*]

print nodes that changed in the indicated time interval

clear

remove all signals from analyzer window

clock [*node|vector* [*val*]]

define a clock bit sequence for the indicated node or vector. This clock signal will then define the actions that occur when either the "**c**" (clock) or "**p**" (phase) commands are issued. If the clock is defined by a vector, then *val* is a binary string equal to the length of the vector. The prefix "**0b**", "**0d**", "**0o**", or "**0h**" may be used to specify the value in (respectively) binary, decimal, octal, or hexadecimal.

d [*node|vector*]...

print node/vector(s) or the entire display list

debug [*options*]

print/set debug state (? for help)

decay [*value*]

set a time limit for undriven nodes to become indeterminate. Otherwise, undriven nodes retain their value prior to becoming undriven, indefinitely. This helps identify problems where dynamic logic is not properly refreshed. A value of **0** (zero) indicates no decay (and is the default value).

display [[-]*option*]

Set or reset display modes. *option* may be one of the keywords **cmdfile**, **automatic**, or **tclproc**. When the keyword is prefixed with "-", the display mode is disabled; otherwise, it is enabled.

cmdfile

When this option is on, commands read from a command file are echoed before they are executed. Default is *off*.

automatic

Print out the current display list (see the "**d**" command) after execution of "**s**" or "**c**" commands. Default is *on*.

tclproc *procedure*

This option declares a Tcl procedure to be executed whenever a node on the display list is displayed. The procedure will be passed, in order, the arguments *nodename*, *bitstring*, and *time*. The procedure must also handle the time signature passed at the end of the watchlist, which takes the same form, with *nodename* = "**time**" and *bitstring* = "**t**". The procedure may choose to ignore any arguments. The Tcl procedure may be removed by passing a null string ("") as the *procedure*.

dump [*file*]

write network state history to file

every *interval* [*start*] { *procedure* }

(command introduced in IRSIM version 9.7.32)

Schedule the command or set of commands "*procedure*" to be executed every *interval* nanoseconds. In the Tcl/Tk version of IRSIM, *procedure* may be any valid set of Tcl commands or IRSIM commands. The procedure is first scheduled to execute at the current time plus *interval*, unless *start* is specified, in which case the procedure will be first scheduled to execute at (absolute) time *start* (nanoseconds). After each execution, the command is rescheduled at *interval* nanoseconds later. In the Tcl version of IRSIM, the **at** command returns a unique ID for the event that can be used for the **cancel** and **get** options (see below). Also see the command **at**. Also see the command **at**.

every get|cancel *id*

(command introduced in IRSIM version 9.7.32)

The **get** option returns the procedure associated with the scheduled event whose identifying (integer) tag is *id*. The **cancel** option cancels the scheduled event whose identifying tag is *id*.

exit [*status*]

exit program with given status (default: 0)

flush (obsolete)

see **histflush**

h *node|vector...*

drive node/vector(s) to 1 (high)

has_coords

print YES if transistor coordinates are available

help [*command*]...

print info on command(s) or available commands

hist [**on|off**]

display or set history collection mode

histflush [*time*] (formerly **flush**)

flush history up to time *time* (default: now). The beginning of the dataset is reset to the indicated time value.

inputs

print currently driven (input) nodes

ires [*time*]

print/set incremental resolution to time

isim [*file*]

read changes from file and incrementally resimulate

l *node|vector...*

drive node/vector(s) to 0 (low)

listnodes

This function returns a list of all the nodes in the (flattened) circuit. IRSIM version 9.7 Revision 4 and newer make use of this to enable a "signal manager" window displaying all traces and allowing traces to be added to or removed from the analyzer display.

listtopnodes [*separator*]

This is a Tcl procedure defined in the Tcl version of IRSIM. It assumes that the simulation netlist was derived from a flattened hierarchical circuit description, and that the node names follow a naming convention of:

instance/instance/.../node

Top-level nodes are those that contain no slash separators. The procedure simply calls the "**listnodes**" function (see above) and returns the subset of names that do not contain slashes. The *separator* is optional; if specified, the given separator character or substring will be used to identify hierarchical nodes in place of the slash character.

listvectors

This function returns a list of all the vectors that have been defined in the simulation.

logfile [[+] *file*]

start/stop log file (+file appends to file)

model [**linear**|**switch**]

print/change simulation model

p

step clock one simulation step (phase). This command requires that a clock has been defined using the "**clock**" command.

path *node|vector...*

critical path for last transition of node(s)

print [*text...*]

print specified text

printp [*n*]

print up to *n* pending events (default: all)

printx

print all undefined (X) nodes

q

terminate input from current stream

query *node|vector*

Tcl version only. Returns the (decimal) value of node *node* or vector *vector*. The returned value is a Tcl object, as opposed to the **assert** short form, which is simply printed to stdout.

querygate *node* (formerly !)

print information about a node's gate connections

querysource *node* (formerly ?)

print information about a node's source and drain connections

R [*n*]

simulate for *n* cycles (default: longest sequence)

readh *file*

read network state history from file

readcver *name*

Read a verilog dump file. This is a Tcl script that makes use of the "**addnode**" command to generate one node per signal, then schedule events as marked in the file, and simulated through those events.

readsim [*prefix*] *name*

Read a .sim format file. IRSIM has the capability to read multiple files and simulate all of them simultaneously. The *prefix* option creates a unique namespace for each file read in, so node names common to two files are not merged in the IRSIM database.

relax [**l**|**h**|**r**]

Attempt to resolve undefined nodes by forcing an initial condition (low logic value) on all undefined nodes. These conditions are applied instantaneously, e.g., like an "l" command, but the logic value state is not persistent. This command usually suffices to resolve nodes that are undefined due to a logical feedback loop. Since this command will obscure truly undefined nodes, it should be used with caution, and only as a last resort. Option "l" forces all undefined nodes to

a low logic value. Option "**h**" forces them to a high logic value, and option "**r**" randomly selects a high or low logic value for each undefined node encountered.

report [*args*]

print/set trace-info or decay report (? for help)

restorestate *filename* (formerly <)

restore the network state from the indicated file

restoreall *filename* (formerly <<)

restor the network state and the value of all inputs from the indicated file.

s [*time*]

simulate for specified time (default: stepsize). *time* is in nanoseconds.

savestate *filename* (formerly >)

write the current network state to the indicated file.

set (obsolete)

see **setvector**

setlog [*file*|**off**]

print/change net-changes log-filename

setpath [[+] *path*]...

print/set/add-to cmd files search path

settle [*value*]

set a minimum time for multiply-driven nodes to become indeterminate. Otherwise, nodes driven by conflicting values become indeterminate based on RC time constant calculations, which will be much too fast, as the condition is primarily controlled by second-order effects. Setting a "settle" value helps prevent propagation of indeterminate values in circuits with pass-transistor logic or back-to-back inverters in delay flip-flops. A typical *value* is normally less than one step size. *value* is a real-valued number, and is interpreted as nanoseconds.

setvector *vector value* (formerly **set**)

Set the value of vector *vector* to *value*. *value* is a binary string equal to the length of the vector. The prefix "**0b**", "**0d**", "**0o**", or "**0h**" may be used to specify the value in (respectively) binary, decimal, octal, or hexadecimal. Also, for backwards compatibility with some IRSIM version of unknown origin, the prefixes "**%b**", "**%d**", "**%o**", and "**%h**" (and equivalently, "**%x**") are also accepted at valid syntax. Negative numbers can be entered as, for example "**-0d1**" and will be translated into a 2's complement binary number with respect to the bit length of the vector being set. For instance, if **A** is a binary vector of length 4 bits, then "**setvector A -0d1**" will set **A** to the binary value **1111**. The notation "**0x**" is allowed for hexadecimal, although be aware that "**0x**" is also a valid binary setting (first bit forced to zero, second bit released from forced input). Thus "**0x0**" is a correct *binary* setting for a vector of size three bits, or a correct *hexadecimal* setting for a vector of any other size.

source

This is the Tcl "source" command, and can be used to read and execute any IRSIM file that contains any mixture of Tcl-compatible IRSIM commands and Tcl/Tk commands. This means that the file must not use the incompatible command names "set" and "flush", which have been replaced by "setvector" and "histflush".

start [-s] *param_file sim_file*

Because IRSIM is a Tcl package, it has an initialization function, which runs automatically upon loading, and a separate function which starts the simulator. This allows IRSIM to be loaded and ready to go before having a ".sim" file ready for input. The **start** command takes the same syntax as the standalone UNIX command prompt **irsim** command.

stats

print event statistics

stepsize [*time*]

print/set simulation step size (nanoseconds).

stop [-] *node|vector*...

pause simulation when node/vector(s) change

t [-] *node|vector*...

start/stop tracing specified node/vector(s)

tcap

print all shorted transistors

time [*command...*]
time given command or program

toggle *node|vector...*
toggle the logical value of a node from 0 to 1 or 1 to 0. For vectors, toggle each bit of the vector independently. Bit values which are undefined are not changed (command introduced in version 9.7.32)

u [*node|vector*]...
drive a node/vector(s) to X (undefined)

unitdelay [*time*]
force transitions to specified time (0 to disable)

until *node|vector* [*mask*] *val count*
simulate until the condition *node* or *vector*, optionally masked with *mask* is equal to value *val*.

V [*node|vector* [*val...*]]
define input sequence for a node or vector. For vectors, *val* is a binary string equal to the length of the vector. The prefix "**0b**", "**0d**", "**0o**", or "**0h**" may be used to specify the value in (respectively) binary, decimal, octal, or hexadecimal.

vector *name node...*
(re)define vector named *name* composed of indicated *node*(s). As an aid to typing long lists of nodes and to aid compatibility with verilog, it is possible to type *node* in verilog vector notation, e.g., **M<1:8>**, under the condition that nodes **M<1>** through **M<8>** exist and that "**M<1:8>**" itself is not a named (single) node. Otherwise, wildcard characters are not allowed in the list of nodes since the order of the nodes cannot be determined. If a vector is redefined, any attributes of the original vector (e.g., being on the display or trace list) are lost.

w [-]*node|vector...*
add/delete node/vector(s) to display-list

when *nodeT valT* { *procedure* }
execute the procedure *procedure* on the next occurrence of the condition that node *nodeT* switches to value *valT*. *valT* may be a concatenated list of types; e.g., "**lh**" will execute on either a rising or falling edge of *nodeT*.

whenever *nodeT valT* { *procedure* }
execute the procedure *procedure* on every occurrence of the condition that node *nodeT* switches to value *valT*. *valT* may be a concatenated list of types; e.g., "**lh**" will execute on both rising and falling edges of *nodeT*. The **whenever** command returns an integer *tag* value that can be used to track the procedure (see below).

whenever get|cancel *id*
The **get** option returns the procedure associated with the scheduled event whose identifying (integer) tag is *id*. The **cancel** option cancels the scheduled event whose identifying tag is *id*.

wnet [*file*]
write network to *file* (smaller/faster than sim file)

x *node|vector...*
make node/vector(s) undriven (non-input)

Xdisplay [*Xserver*]
print/set X display (for analyzer)

Analyzer GUI Interface Commands

The original "analyzer" graphics display was not associated with any command-line interface commands. This situation was rectified, by necessity, to enable the use of the analyzer and GUI under Tk. The commands are as follows:

base get|set [*trace*] [*base_type*]
Set or query the numerical base for printing vector values on the trace. Valid *base_type* names are: **binary**, **octal**, **decimal** (unsigned decimal), **signed** (signed decimal), and **hexidecimal**. Signed values are in 2's complement notation with respect to the length of the vector. Values other than binary are valid for vector lengths up to 64 bits.

marker [**1|2**] *option...*
Set or change the position of the two vertical markers. Valid options are:

get

Get the marker position

set *trace time*

Set the marker to mark signal *trace* at time *time*

move *time*

Move the marker to time *time*

off

Turn off the selected marker

print *option...*

Handle printing options, which are as follows:

banner [**on**|**off**]

If true, print the title (see below) as a header on the output page.

file *filename*

Print a PostScript version of the screen to the indicated file named *filename*.

legend [**on**|**off**]

If true, print a legend of each trace on a second output page (off by default).

outline [**on**|**off**]

If true, print an outline box around the output graph.

times [**on**|**off**]

If true, print simulation times along the X axis of the output graph.

title *string*

Set the title for the page banner to the string *string*. This also changes the title of the analyzer window, displayed in the upper left-hand corner.

simtime *option...*

Report or set the time value for various points in the analyzer window. Valid options are:

begin

Report the time value of the start of the dataset. Normally this is zero, unless the **histflush** command has been issued.

end

Report the time value of the end of the dataset, which is the current simulation time.

left [*value*]

Return the time value represented by the left edge of the analyzer window. If *value* is given, then adjust the analyzer window view so that the left edge is at the indicated time.

right [*value*]

Return the time value represented by the right edge of the analyzer window. If *value* is given, then adjust the analyzer window view so that the right edge is at the indicated time.

delta

Return the time span represented by the distance between the edges of a split marker.

marker

Return the time represented by the horizontal position of the vertical marker (if it is active).

cursor

Return the time represented by the position of the mouse pointer in the analyzer display window.

move [**+**|**-**]*value*

Move the analyzer display view so that the value *value* is positioned at the left edge of the window. If a plus or minus symbol is given before *value*, then *value* is treated as a time relative to the current time at the left edge of the window.

scroll [**on**|**off**]

Set or report the status of auto-scroll. When auto-scroll is on, the analyzer display updates every simulation timestep to set the current simulation time at the right edge of the display window. Note that even when auto-scroll is active, if the current timestep is outside of the window display prior to the simulation step, then the window view will not be updated.

trace *option...*

Add, modify, or remove a trace from the analyzer.

Valid options are as follows:

top *name*

Return the pixel Y value of the top edge of the indicated trace in the analyzer display window.

bottom *name*

Return the pixel Y value of the bottom edge of the indicated trace in the analyzer display window.

order *name*

Return the order (numbered from zero, the trace at the top of the window) of the indicated trace.

base [*name*] [*value*]

Return or set (if *value* is given) the numeric base used to display the value of a vector. Known values are **binary**, **octal**, **decimal**, and **hexidecimal**.

class *name*

Return the class (which is either "node" or "vector") of the indicated trace.

cursor [*value*]

Return the name of the trace which is closest to the pixel Y position *value*. If *value* is not given, then return the name of the trace that is closest to the mouse pointer position. If no trace is near the indicated value or pointer position, then nothing is returned.

input *name*

Indicate whether the value of *name* (or each bit, individually, if *name* is a vector) is a driven (input) signal or not. Driven bits are returned as value "i", and undriven bits are returned as value "-".

list all|nodes|vectors

Return a list of traces displayed in the analyzer. This can be either a list of all traces, or it can be restricted to list only those traces that are nodes or those that are vectors.

select [*name*]

Set the currently selected trace to be the trace indicated by *name*. If no trace name is given, then return the name of the currently selected trace, if any.

value *name*

Return the instantaneous value of the trace named *name* at the time represented by the vertical marker. If the marker is not active, then nothing is returned.

bits *name*

Return the number of bits in the trace indicated by *name*. If *name* is a node and not a vector, then the result is always 1.

remove *name*

Remove the trace named *name* from the analyzer display.

characters [*value*]

Limit the number of characters displayed for each trace name on the left side of the display to the integer number *value*. This prevents long hierarchical names from dominating the analyzer display window. If no *value* is given, then the current limit value is returned. The default character limit is 15.

move *name1 name2*

Swap the relative positions of traces *name1* and *name2* in the analyzer window.

zoom [*in|out*]

Change the time span of the window view, effectively zooming the time axis in or out.

Statistics Package Commands

If the "stats" packages has been compiled in (normally true), then the following additional commands are available:

histev [*clear|off|on*]

enable/disable/clear event activity record

evstats [*file*]

Connection-List Statistics Package Commands

If the "cl statistics" packages has been compiled in (normally true), then the following additional command are available:

clstats [*file*]
print connection-list statistics

Fault Simulator Package Commands

If the "fault simulator" package has been compiled in (normally true), then the following additional command is available:

faultsim *infile* [*outfile*]
do stuck-at fault simulation

Power Estimation Package Commands

If the "power estimation" package has been compiled in (normally true), then the following additional commands are available:

powlogfile [[+] *file*]
start/stop power logfile (+*file* appends to file). Typically, this is called as "**powlogfile /dev/null**" to redirect output, as the power summary and histograms are more useful than the per-step output.

powtrace [-] *node|vector...*
start/stop power tracing specified node/vector(s). The *node* string can be wild-carded. Typically, this is called as "**powtrace ***" to enable power calculations on all nodes.

sumcap
print out sum of capacitances of all nodes.

power *node*
specify that the node named *node* is a power supply node. All power and ground nets must be specified for the power estimation to work correctly. Note that specifying a net as a power supply is not sufficient to set the network as logic high. The net must be specifically set to logic high with the command "**h** *node*". Multiple nodes can be declared power supplies.

ground *node*
specify that the node named *node* is a ground return node. All power and ground nets must be specified for the power estimation to work correctly. Note that specifying a net as a ground return is not sufficient to set the network as logic low. The net must be specifically set to logic low with the command "**l** *node*".

vsupply [*net*] [*value*]
Set supply voltage = *value* Volts. The *net* gives the name of the power supply net. Multiple power supplies with different values can be specified.

powstep
Toggle display of power estimate for each step. Note that this command is mutually exclusive with **powhist**.

powhist init *min max* [*buckets*]
Initialize power calculation histogramming. The histogram method is mutually exclusive with **powstep**. The histogram data is taken over *buckets* bins with the first bin at power value *min* and the last bin at power value *max*. Values for *min* and *max* are given in milliwatts.

powhist capture|print|reset
Once power histogramming is enabled, data are analyzed using the following command options:

powhist capture
IRSIM accumulates power data (total switching energy) on every device every time the device switches, resulting in a total value of energy expended in the system.

The **capture** option divides the energy by the time over which data have been accumulated, arriving at a power value (total energy divided by the time interval). The resulting power value is histogrammed according to the histogram values specified in the **powhist init** command option, and the energy values are reset to zero. Typically, this command option is called once per clock cycle, so a typical setup would be to declare "**every time {powhist capture}**" and set *time* to the clock period, resulting in a histogram of clock cycle power.

powhist print

This command option prints the histogram data, with one histogram bin per line, with three entries being the bin starting value, the bin ending value, and the number of entries within the bin.

powhist reset

This command resets the histogram data and terminates the collection and histogramming of power data.

powquery

This command returns the current calculated value of the total power, which is the total accumulated energy divided by the time interval since the use of **powstep**, **powhist init**, or **powhist capture**, for the nodes specified by the **powtrace** command.

User Subcircuit Package

IRSIM can implement arbitrary circuit components, not just the internally-defined transistors, resistors, and capacitors. Subcircuits can be defined and used in the ".sim" file with the syntax:

```
x input1... [output1]... [options] subcircuit_name
```

This line defines a subcircuit named *subcircuit_name* that connects to a list of input and output nodes starting with *input1*. The remaining arguments *options* may be anything, although it is recommended that they be restricted to "*key=value*" pairs defining parameters used by the subcircuit.

There are three different ways to define user subcircuits in IRSIM. Two of them involve defining the subcircuits in precompiled C code. This requires that C routines be created for all subcircuits needed.

The first of the two pre-compiled methods is the original one, and results in a special version of the IRSIM executable that understands these definitions. For the non-Tcl-based IRSIM, this is the only choice for making user-defined subcircuits. See the source distribution's "usersubckt" subdirectory for examples of code defining compiled-in subcircuits.

All subcircuit definition methods are available with the Tcl-based version of IRSIM, although the method mentioned above is deprecated in favor of the second method. In the second method, a file of C routines defining a library of subcircuit components is required, and it is compiled as a shared-object library. The shared-object library can then be loaded at run-time. This method makes it possible to maintain a set of different libraries for different projects. Note that the shared-object library must be loaded prior to any .sim-format file containing instances of the subcircuits defined in that library. A convenient way to handle this is *not* to read the .sim file from the OS command line, but to load the shared-object library and load the .sim file (in that order) from either the command file script or the IRSIM command line, after startup:

```
# Tcl command file script for IRSIM load /usr/local/lib/irsim/tcl/diglib.so
readsim mysimfile.sim
```

The `tclsubckt` subdirectory of the IRSIM source has an example file, `diglib.c`, containing examples of digital standard cells defined C code routines. Each subcircuit requires an evaluation function, and optionally an initialization function.

The third method of defining a subcircuit is slower but does not require any compilation. To do this, simply define all subcircuits as procedures in Tcl at startup. This can be done in a similar way to the manner mentioned above, in which IRSIM is started without loading the .sim file, then the command file script defines all the necessary subcircuits before loading the .sim file with the "**readsim**" command.

Each subcircuit requires two procedures, an initialization procedure and an evaluation procedure. The syntax of these procedures is defined below:

```
proc subcircuit_name_init {} {  
  procedure_body  
}  
  
proc subcircuit_name_eval { instr outstr udata } {  
  procedure_body  
}
```

The initialization procedure takes no arguments. It must return a list of values describing the circuit. This list consists of, in order:

num_inputs

An integer declaring the number of input nodes in the subcircuit. Every circuit must have at least one input node, even if it is just a "dummy" node used to force evaluation of the subcircuit. A subcircuit is only evaluated when one of its inputs changes state.

num_outputs

An integer declaring the number of output nodes in the subcircuit. A subcircuit may legally have zero outputs (the subcircuit may only be producing diagnostic output, or communicating with another tool).

res_list

A list of (*num_outputs* * 2) real-valued numbers. The first half are the values of the pull-up transistor "on" resistance of the subcircuit's output driver pFET, one value per output. The second half are the values of the pull-down transistor "on" resistance of the subcircuit's output driver nFET, one value per output. All values are in Ohms. If the circuit has no outputs, this term should be an empty list ("{}").

user_data

This term is optional and may contain any data whatsoever. *user_data* is a single value, however, so multiple values in the data should be collected together into a list. A unique instance of this value is saved for each instantiated subcircuit. It is passed directly to the evaluator procedure, which may modify it. Since the user data is a Tcl list, the evaluator procedure may modify the list directly using Tcl commands like "**lreplace**". This allows each subcircuit instance to maintain internal states. If an "**x**" record in the .sim file has parameters between the last I/O node and the subcircuit name, those parameters will be appended to the *user_data* list for the instance, and should be handled by the evaluator procedure (see below).

The evaluator procedure will be passed the two strings "*instr*" and "*outstr*", which are the values of the input and output pins, in character form ("x", "0", and "1"). If the initialization procedure returned a non-NULL value for *user_data*, and/or the instance record in the .sim file contained extra parameters, then this value will also be passed to the evaluator procedure. The evaluator procedure must return a list containing the following values:

newoutstr

This is a string containing character values for each of the output nodes. The characters may be from the set "0", "x", "z", and "1", and also "l" and "h", case insensitive.

delays

This is a list of real-valued numbers representing the delay to each of the outputs. Note that if the outputs have different delays relative to different inputs, it will be necessary for the subcircuit to maintain an internal representation of the previous state of all inputs, and determine which delay value is appropriate. Also note that the output nodes are given no capacitance, so the delay value should include any transition delay in the outputs due to charging their own capacitance.

For example, the procedure defining an inverter might look like:

```
proc inverter_init {} {  
  # Inverter is 1 input, 1 output, 500 ohms "on" resistance  
  return {1 1 {500.0 500.0}}  
}
```

```

proc inverter_eval {in out ulist} {
    # Inverter delay is 10ps. ulist is not used.
    switch -- $in {
        0 {set rval 1}
        1 {set rval 0}
        default {set rval x}
    }
    return {$rval 10}
}

```

Procedures for most combinatorial logic are conveniently implemented as logic tables using the Tcl **switch** statement, particularly simple when using the **-glob** option with wildcard characters. The return delay value is given in picoseconds. There is no limit on the number of subcircuit procedures that may be defined.

Important note: The .sim file syntax for subcircuits shown above has been changed from the original IRSIM user subcircuit definition, and is *not* backwardly-compatible. Old user-defined subcircuits need to be rewritten to accomodate the new syntax by removing the "number of inputs" and "delay" parameters from the .sim file and putting them in the subcircuit definition tables in the source code.

Random Number Generator Package Commands

The Tcl/Tk-based version of IRSIM comes with a separate Tcl package extension that creates the Tcl command "**random**", which is based on the C stdlib "rand48" routines.

The package itself is general-purpose and can be run independently of IRSIM. In any "wish" or "tcltk" interpreter environment, run the command "**load /usr/local/lib/irsim/tcl/random.so**" to load and initialize the package. Within IRSIM, this package is loaded automatically. Within IRSIM, the most helpful use of the package is the "**random -bitstream *n***" option, to generate uniformly random input vectors into a circuit. For example:

setvector A [random -bitstream 16]

will set the vector **A** to a random value of 16 bits.

The package creates a single new command "**random**" whose options are as follows:

no arguments

Returns a number taken uniformly from [0..1)

-reset

Reseeds the generator using a combination of the current process ID and current time.

-seed *n*

Reseeds the generator with the specific seed value *n*.

-integer [*a* [*b*]]

Generates a uniformly-distributed random number in the indicated range (see below), and returns the value rounded down to the largest integer less than or equal to the floating-point value. By default, the number returned is within the full range of a 32-bit signed integer. With option "**random -integer *a***", the number returned will be between *-a* and *+a*. With option "**random -integer *a* *b***", the number returned will be between *a* and *b*. The routine is capable of returning numbers up to 64 bits unsigned.

-bitstream [*n*]

will generate a uniform random binary value of length *n*. The value returned is in string form, with effectively unlimited length. By default, the number returned is a 32-bit value (32-character string).

-normal *m* *s*

Returns a floating-point value taken from a gaussian distribution with mean *m* and standard deviation *s*.

-exponential *m*

Returns a floating-point value taken from an exponential distribution with mean *m*.

-uniform *a* *b*

Returns a floating-point value taken from uniform distribution on [*a*,*b*).

-chi2 *n*

Returns a floating-point value taken from chi2 distribution with *n* degrees of freedom.

-select *n list*

Returns *n* elements to be selected at random from the Tcl list *list*, with replacement.

-choose *n list*

will cause *n* elements to be selected at random from the Tcl list *list* without replacement.

-permutation *n*


will return a permutation of 0..*n*-1

-permute *list*

will return a permutation of Tcl list *list*

The Analyzer GUI Interface

The analyzer GUI window displays node and vector values as a graph over time. Traces are added using the **"ana"** command and manipulated or removed using the **"trace"** command (see command descriptions, above). The main display window consists of three parts: A list of node names on the left, the traces in the center, and a list of instantaneous values on the right.



Screen shot of the analyzer graphic display in Tcl/Tk-based IRSIM 9.7

Both individual nodes and vectors can be plotted. Individual nodes show values low, high, and indeterminate (with a crosshatch pattern). Vectors display a view showing regions split at points where the value of any signal in the vector changes. Where space allows, the value of the vector is printed in the chosen base.

To manipulate a signal, click on the name of the signal on the left. The button will turn orange, indicating that the signal has been selected. When selected, use of the **"Base"** menu will change the numeric base in which that signal is displayed. All vectors default to binary. When any other base is chosen, any unknown bit will cause the corresponding digit to be displayed as "X" (in octal or hexadecimal) or as "?" (in decimal, which has no one-to-one bit-to-digit correspondence).

The analyzer window features a vertical marker. The marker is activated by clicking the left mouse button in the main trace display window, and disabled by clicking the middle mouse button. It can be moved by dragging while holding down the left mouse button. While active, the time position marked by the marker is displayed over the main trace window, and the instantaneous value of each trace in the display at the time represented by the marker is displayed on the right-hand side of the window. Clicking and dragging while holding the right mouse button down shows a split marker, and displays the time at both ends of the split marker, as well as the time span between them. If the **"z"** key is pressed while displaying a split marker, the view will expand to make the area inside the split marker the width of the full screen.

At the bottom of the window, there is a scrollbar and scroll buttons. The solid part of the scrollbar represents the size and position of the window view relative to the simulation database. The whole scrollbar represents the database from time zero to the current time. The scrollbar can be dragged left or right with the left mouse button to change the time position of the screen. The ends of the scrollbar can be repositioned by dragging while the right mouse button is held down, which changes the zoom factor. The circle icon at the lower left causes the window view to display all recorded data from time zero at the left edge to the current time at the right edge. The double-left arrow positions time zero at the left edge (without changing the zoom); the double-right arrow positions the current simulation time at the right edge of the window (without changing the zoom factor). The single right and left arrows scroll the view forwards and backwards in time, respectively.

At the top of the window, the "Zoom" button changes the zoom factor (only one of several methods to change the zoom factor; see the paragraph above). The "Base" button changes the displayed numeric base of the selected trace. The "Window" button pops down a menu that has a trace manager popup window, that allows quick access to all of the nodes in the simulation, and defined vectors, and allows the user to choose which ones to display in the analyzer. The "Print" button allows one to generate a PostScript file of the contents of the trace window, with various options controlling what is added to the printout.

Underneath the menu bar, two or three values are displayed. The value on the left is the time represented by the left edge of the trace display window. The value on the right is the time represented by the right edge of the trace display window. The value in the center is only shown when the vertical marker is active, and shows the time represented by the marker position.

The analyzer can also be manipulated from the command line. In addition to the command **"analyzer"** that generates the window and adds traces, its behavior is controlled by the five commands: **"base"**, **"marker"**, **"print"**, **"simtime"**, and **"trace"** (see list of commands, above, for a detailed description).

Using IRSIM With Tcl/Tk-Based Magic

(See a [screenshot](#) of IRSIM 9.6 running under magic 7.2 with Tcl/Tk (this is a bit out of date and needs to be updated to IRSIM 9.7 with the Tk-based analyzer window). The layout is the standard magic tutorial **tut11a.mag** and the setup was made with a slightly modified IRSIM command file [tut11a.cmd](#), executed using the Tcl command **source**).

The IRSIM digital switch simulator was a casualty of the switch to the Tcl/Tk interpreter, until Magic 7.2 revision 14.

The focus of the Tcl/Tk version of Magic is to take advantage of the interpreter's ability to load packages on the fly. The result is that disparate programs can run at the same time. Any program can call routines in any other program, although this is normally done via evaluation of interpreter commands.

To facilitate having other programs pass information back to magic for display, a new magic command "element" was added. This general-purpose method lets other programs print text, lines, and boxes on top of the magic layout.

To make use of IRSIM under Tcl/Tk-based Magic, it is necessary to have IRSIM version 9.6 Revision 1 or newer. IRSIM must be compiled with the Tcl interpreter (the default configuration).

IRSIM works in the following manner. IRSIM is, like magic, exttosim, and exttospace, compiled as a Tcl package, namely as a shared-object library (.so) file. There is a standalone script "irsim" which is normally in **/usr/local/bin**. Invoked from the script, IRSIM runs under Tcl by itself, and can simulate from any ".sim" file, but cannot interact with magic layout. However, there is also a command "irsim" in magic. Invoked as "irsim", IRSIM loads as an additional package in the Tcl interpreter which launched magic. Able to detect the "magic" package in Tcl, IRSIM then adds a group of commands designed for interaction between magic and IRSIM. The "irsim" procedure is smart enough to go looking for a ".sim" file matching the magic layout in the window, and generate one automatically if it doesn't exist. Because the IRSIM and magic packages share space in the Tcl interpreter, commands for both are

available at any time from the command line. There is no switching between modes to access the simulator. The only caveat is that some IRSIM commands conflict with magic or built-in Tcl/Tk commands. These commands require a prefix "irsim::". For example, "clock" is a command used by Tcl to time the execution of procedures. To get the "clock" command in IRSIM, it is necessary to use "irsim::clock". Several IRSIM commands conflict with Tcl syntax, and so have been renamed. These are indicated below.

The IRSIM/Magic procedures defined when running IRSIM under Magic are as follows:

irsim [*param_file* [*cellname*]]

This command runs **irsim::start** in a somewhat intelligent manner. If *cellname* is not supplied, it defaults to null. Under this condition, the toplevel cell in the magic layout window is assumed to be the one to be simulated. If *param_file* is missing, it will be substituted with the default value "scmos100.prm", corresponding to 1.0 micron technology.

watchnode [*node_name* [*magic_color*]]

This command causes IRSIM to display node values on top of the magic layout. With no arguments, "watchnode" expects a node to be (uniquely) selected in magic. It uses the magic "getnode" command to get the name of the node and map it to an IRSIM node name. The current value of the node is then printed on the layout, centered on the select box, and updated as needed by the simulation. If *node_name* is specified, IRSIM uses the magic "goto" command to find some valid label or paint in the layout belonging to the node, and displays the value there. *magic_color* is an optional color name from the "long names" column in the magic ".dstyle5" display styles file. If not specified, it defaults to "white".

Note that *node_name* may also be an IRSIM vector name, in which case the magic "goto" command fails and the value is placed at the position of the magic cursor box.

unwatchnode *node_name*

This erases the simulated node *node_name* from the magic layout.

movenode *node_name*

This command recenters the label for simulated node *node_name* to the position of the magic cursor box. This is useful if the value chosen by "watchnode" is in a poor position for viewing.

watchtime


This writes the simulation time in nanoseconds on top of the magic layout, placed at the position of the magic cursor box.

unwatchtime

Erases the simulation time display from the magic layout.

movetime

Recenters the simulation time display on the magic cursor box position.



Results of "watchnode" and "watchtime" on the magic layout.

IRSIM defines a callback function which returns values for all watched nodes and vectors on every update. This function is invoked as an extension of the IRSIM "display" command as follows:

display tclproc *procedure_name*

The Tcl procedure *procedure_name* must be defined as a procedure taking exactly three arguments:

```
proc procedure_name { name value tval } {  
... (procedure body) ...  
}
```

IRSIM will call *procedure_name* once for each variable needing a display update, and fill the values for **name**, **value**, and **tval**. Argument **name** is the name of the variable (node or vector) to be updated. Argument **value** is the value of the variable (one character for nodes, or a character string for vectors). Argument **tval** is the time of the update. In order to facilitate the batch processing of all nodes and vectors which are updated at the same time value, *procedure_name* will be called once more after all variables have been updated, with the syntax

procedure_name **time t tval**

The procedure *procedure_name* should check for the condition "**#{value}**" == "**t**" to differentiate this call from calls on node and vector names.

Technology (.prm) files for IRSIM version 9.7:

Parameter file syntax

An IRSIM parameter (".prm") file consists of a series of lines, each of which begins with a keyword that determines how the rest of the line is interpreted. Each line is of the form:

parameter_name value1 ... valuen

A ';' character indicates that the rest of the line is to be ignored (a comment). The following parameters are understood by IRSIM:

lambda *real_number*

Multiplicative constant to convert from .sim file linear dimensions (e.g., lambda) to micrometers.

capm2a *real_number*

2nd metal capacitance per unit of area (in pF/sq-micron).

capm2p *real_number*

2nd metal capacitance per unit of perimeter (in pF/micron).

capma *real_number*

1st metal capacitance per unit of area (in pF/sq-micron).

capmp *real_number*

1st metal capacitance per unit of perimeter (in pF/micron).

cappa *real_number*

poly capacitance per unit of area (in pF/sq-micron).

cappp *real_number*

poly capacitance per unit of perimeter (in pF/micron).

capda *real_number*

n-diffusion capacitance per unit of area (in pF/sq-micron).

capdp *real_number*

n-diffusion capacitance per unit of perimeter (in pF/micron).

cappda *real_number*

p-diffusion capacitance per unit of area (in pF/sq-micron).

cappdp *real_number*

p-diffusion capacitance per unit of perimeter (in pF/micron).

capga *real_number*

MOSFET gate capacitance per unit of area (in pF/sq-micron).

cntpullup *boolean*

If **0**, then the gate capacitance of depletion transistors used as pullup devices is NOT included in the output node. This parameter is only significant for nMOS designs which have depletion

devices. Any number other than 0 will set this flag, this is true for all other 'boolean' parameters below.

diffperim *boolean*

If not **0** then the sidewall diffusion capacitance is subtracted from the source and drain nodes. The capacitance is calculated as (*capdp* * *width*). This applies to both n and p transistors.

subparea *boolean*

If not **0** then the poly over transistor capacitance is subtracted from the gate node. The capacitance is calculated as (*cappa* * *area*).

diffext *real_number*

Each transistor is assumed to have a rectangular source and drain diffusion extending *real_number* units wide and *transistor-width* units high. The effect of the diffusion extension is to add some capacitance to the source and drain nodes of each transistor. This is useful when processing the output of NET to improve the capacitive loading approximations without adding explicit load capacitors and is not very useful for other types of networks and will only be applied if *real_number* is different from **0**. *diffext* is specified in microns (contrary to what some .prm files may say).

lowthresh *real_number*

Normalized logic low voltage threshold. Any voltage \leq *lowthresh* is assigned a low logic value (**0**).

highthresh *real_number*

Normalized logic high voltage threshold. Any voltage \geq *highthresh* is assigned a high logic value (**1**). Setting *highthresh* = *lowthresh* will disallow the unknown (**X**) node value. Note that setting *highthresh* lower than *lowthresh* will not result in nodes having hysteresis and will probably result in disaster. Numbers commonly used for the thresholds are **0.4** and **0.6**.

device *device_type device_name [value]*

This line declares that the name *device_name* is a subcircuit in an input .sim file (i.e., a x record). When encountered in an input file, IRSIM should parse the record as if the device is a device of type *device_type* with arguments in the .sim file being specific to the device type, as follows:

nfet

The device is an nFET transistor and has a device line of the syntax

x g s d b l=l w=w x=x y=x device_name

When declaring a device type **nfet**, the *value* record is not used. Transistor values are given by the *resistance* statement (see below). Note that this differs from the **n** syntax in the .sim file by specifying a bulk terminal connection.

pfet

The device is a pFET transistor and has a device line of the syntax

x g s d b l=l w=w x=x y=x device_name

where *g* is the gate node, *s* is the source node, *d* is the drain node, and *b* is the bulk node. When declaring a device type **pfet**, the *value* record is not used. Transistor values are given by the *resistance* statement (see below). Note that this differs from the **n** syntax in the .sim file by specifying a bulk terminal connection.

resistor

The device is a resistor and has a device line of the syntax

x d0 t1 t2 l=l w=w x=x y=x device_name

where *d0* is the (unused) device node, and *t1* and *t2* are the resistor terminals. The *value* in the statement is the resistance in ohms per square and tells IRSIM how to convert the device length and width into a resistance.

capacitor

The device is a capacitor and has a device line of the syntax

x top bot l=l w=w x=x y=x device_name

where *top* is the capacitor top node and *bot* is the capacitor bottom node. The *value* in the statement is the capacitance in picoFarads per lambda squared (see **lambda**, above) and tells IRSIM how to convert the device length and width into a capacitance.

diode

The device is a diode and has a device line of the syntax

x plus minus l=l w=w x=x y=x device_name

where *plus* is the diode's positive terminal and *minus* is the diode's negative terminal.

The *value* in the statement is the capacitance of the diode's junction area in picoFarads per

lambda squared (see **lambda**, above). Diodes are not directly simulated in IRSIM but will be converted to an equivalent capacitance that affects timing.

resistance *type context width length resistance*

Defines the resistance (in ohms) for a transistor of type *type* having *width* - *length* dimensions when connected as in *context*. Note that the dimensions are in microns (not lambda). All of *width*, *length*, and *resistance* are real numbers.

Values for *type* may be one of the following:

n-channel

Corresponds to an **n** or **e** device in a .sim file. These resistances will also be used as default values for any **nfet** device that is not given its own resistance values.

p-channel

Corresponds to a **p** device in a .sim file. These resistances will also be used as default values for any **pfet** device that is not given its own resistance values.

depletion

Corresponds to a **d** device in a .sim file.

resistor

Corresponds to an **r** device in a .sim file (not to be confused with an **R** device, which is a lumped resistance).

device_name

Corresponds to an **x** device in a .sim file. The *device_name* must have been previously declared in a **device** statement (see above).

Values for *context* may be one of the following:

- **dynamic-low**
- **dynamic-high**
- **static**
- **power** (*parsed but not used*)

See the section "Generating the resistance tables" below for the explanation of these resistance types. In addition, the *context* can be qualified with "**-with-drop**" but IRSIM does not use this context at all. It is only parsed for backward compatibility with old parameter files.

Also for the sake of backward compatibility, the static context is not a resistance altogether, but it is used to approximate the transistor's transconductance (gm) at the switching point (between the linear and saturation regions of operation).

For a CMOS process, it is sufficient to specify the following:

n-channel dynamic-low

n-channel dynamic-high

n-channel static

p-channel dynamic-low

p-channel dynamic-high

p-channel static

You can specify a series of resistance values as IRSIM will use linear interpolation for any non-listed value. This may be particularly important for technologies with channel lengths less than approximately 2um, since short channel effects become important and transistor resistance does no longer scale with channel length.

The remaining parameters that are absolutely necessary are

lambda

capga

The value of **lambda** is used to scale all distance values in the input ".sim" file. Any value that is not qualified with a metric prefix (e.g., "n" or "u" for nanometers or micrometers) is assumed to be in integer multiples of lambda. Traditionally, **lambda** is equal to half the length of a minimum-length transistor gate. However, in practice it may be almost anything.

The value of **capga** is the most important of all, because virtually every process model for MOS transistors includes the gate capacitance as part of the model, and it is rare to find any CAD tool that extracts gate capacitance for a device, since one does not want to count the gate capacitance twice. The ".sim" file format is much simpler than a SPICE deck and lacks all detailed model information. The simplified model information is contained in the parameter file, and it is critical that gate capacitance be accounted for; otherwise, the capacitive load on the input of every transistor will be missing.

Likewise, most CAD extraction tools will extract all of the other capacitance values mentioned above. Magic's `ext2sim` command will generate a *lumped capacitance* per node, containing the sum of all capacitances between each mask layer and the substrate (ground) for each layer type directly connected to that node. This means that the remaining values are unnecessary for the simulation of extracted layout. However, simulation files generated by schematic capture will not contain the detailed parasitic capacitance information per node. Because transistor source and drain parasitic capacitances are significant, often the schematic capture tool will generate area and perimeter information about the source and drain area for each transistor, an estimate based on the gate width and a standard minimum-size layout for the transistor. If the schematic capture tool provides this information as "s=A_...,P_..." and "d=A_...,P_..." entries for each transistor, then these area estimates are accounted for. In this case, you will want to have valid entries for parameters **capda**, **capdp**, **cappda**, and **cappdp**, and a valid entry for **lambda**, since source and drain area and perimeter attributes are given only in lambda (and lambda-squared) units.

Finally, if you have a .sim file generated by schematic capture and which does not contain source and drain area and perimeter information, IRSIM can make a first-order estimate of source and drain capacitance based on the given gate width, and using the value **diffext** as the "standard" width of a source or drain diffusion region (in microns). The simulation result should not be construed as having accurate timing information, but will be much closer to a realistic value than having no source and drain capacitance at all.

Notes on parameter files for .sim files created by ext2sim (magic)

By far the most common (and probably the only remaining) source of .sim files is from the EDA layout tool **magic**. Here is a summary of what you can expect to find in the .sim netlist files it produces:

lambda:

The first line of a .sim file created by ext2sim contains the following:

```
| units: scale tech: tech_name
```

where *tech_name* identifies the technology and *scale* is the multiplicative constant that converts .sim dimensions to (usually) centimicrons. The correct lambda value for these files should then be *scale*/100.0. Note that IRSIM ignores this line in the .sim file so it is important that it be specified correctly in the .prm file. Magic technology files in magic versions 7.0 and higher may specify output dimensions to be in nanometers or even in angstroms, so if you're not sure, it is a good idea to check the length value of a minimum-length transistor and confirm that multiplying that by **lambda** results in the correct length, in microns.

cap*:

Magic generates the appropriate layer capacitances when extracting a circuit and IRSIM will not make use of any the following: capm2a, capm2p, capma, capmp, cappa, cappp, capda, capdp, cappda, cappdp. These parameters can be safely set to zero. Since their value will not be used, you can include the appropriate values for documentation.

capga:

Capga should always be specified (unless you want to ignore gate capacitive loading - probably a bad idea). Capga can be calculated directly from the process parameters (or corresponding spice deck model) as follows:

$$capga = Cox / Tox$$

where:

$$Cox = (\epsilon_0)(\epsilon_{ox}) \approx (8.85)(3.9) \text{ pF/m} = 34.5 \times 10^{-6} \text{ pF/um}$$

so that:

$$capga = (34.5 \times 10^{-6}) / Tox \quad (\text{note: } Tox \text{ in um}).$$


Note: There should really be a $capnga$ & $cappga$ for n and p FETs, but this is not a severe problem yet. Some processes specify an *effective* oxide thickness Tox , and this effective thickness is different for the two types of complementary FETs, even when the oxide itself is the same thickness throughout.

resistance:

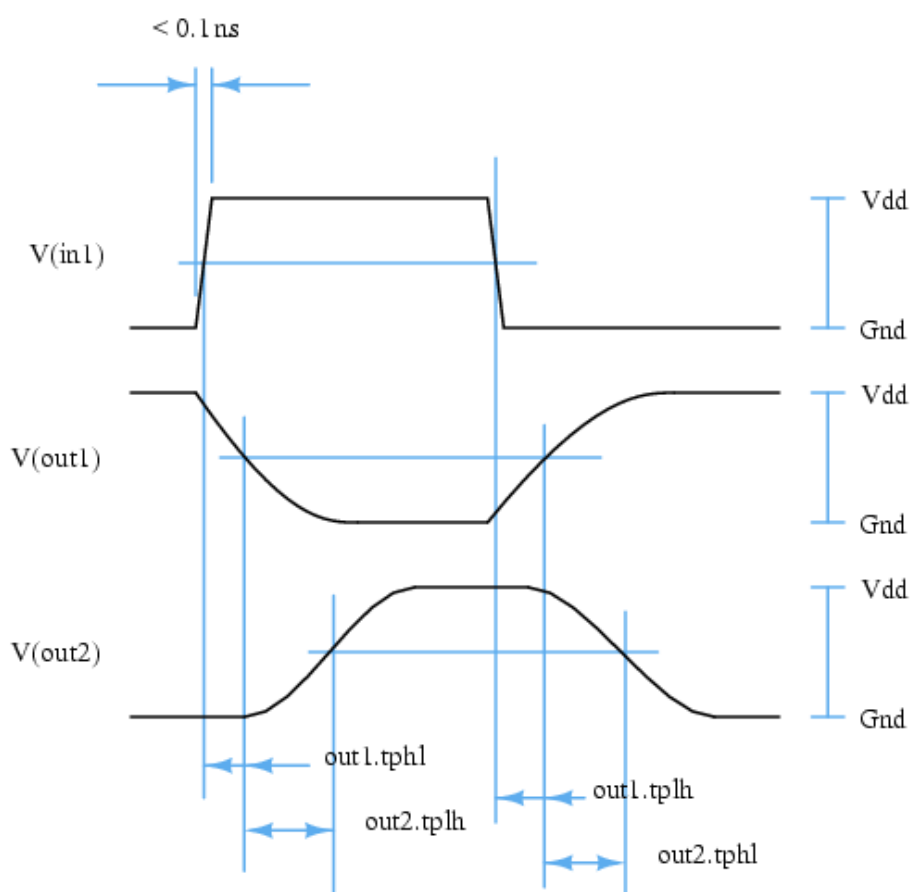
Magic never generates the pullup context and, if a depletion device is connected as a pullup, IRSIM will recognize it as such. There is no need to specify that context.

Generating the resistance tables

If you have a SPICE model for the transistors in the process, you can use SPICE (or some other circuit simulator) to approximate the resistances. The following circuit (for a CMOS process) can be used to extract the resistances used by IRSIM.



Apply a pulse with a rise/fall time of $\leq .1\text{ns}$ to $in1$ and measure the delay from $in1$ to $out1$ and from $out1$ to $out2$ for both transitions. Make sure that Cap is sufficiently large (1pF is reasonable), otherwise the delay will be determined primarily by the internal (parasitic) capacitance of the driving device. It is a good idea to ratio the 2 inverters so that the midpoint lies in the middle and you can easily measure the delay. Make sure the pulse is sufficiently long to allow the signals to reach a steady state.



From the delays measured, you can extract the resistances as follows:

```
n-channel dynamic-low = out1.tphl / Cap
p-channel dynamic-high = out1.tplh / Cap
n-channel static = (out2.tphl^2 - out1.tphl^2) / (out1.tplh * Cap)
p-channel static = (out2.tplh^2 - out1.tplh^2) / (out1.tphl * Cap)
```

note: $a^2 \Rightarrow a * a$

You will see from the figure that the primary (first-order) delays are due to the effective channel resistance (the dynamic-high and dynamic-low values) when the MOSFET is driven far into saturation by a step function input. As the source and drain parasitic capacitances are negligible, the rise time $t = RC$. As mentioned further up this page, the static "resistance" values are not really resistances at all, but estimate the gm of the transistor. They represent a second-order correction on the simple RC time constant that takes into account the much higher resistance of the transistor in the linear region, before it reaches saturation. This region cannot be seen when the transistor is driven by a hard step function, which is why the static resistance is estimated by comparing the difference between the response to a hard step function and the response to the finite rise or fall time of the stage in front of it.

You can similarly measure the delay for an nFET driving a large capacitance high to determine "n-channel dynamic-high" and a pFET driving a large capacitance low to determine "p-channel dynamic-low". This setup requires an nFET transistor connected to power and a pFET transistor connected to ground, a circuit configuration that is rarely used except in the occasional hysteretic inverter. It is therefore unlikely that accurate values would ever be needed for these two parameters.

By the way, I have not worked out the derivation of the static values shown above, although I assume they are valid. If anyone has a derivation, please pass it on and I will include it on the web page.

Most designs (especially those that use digital standard-cell libraries) use exclusively minimum length devices, so it is usually sufficient to specify the above resistances for a single device length only. IRSIM will then calculate the resistance values for each device based on the device width and length and a single value of resistance in ohms per square. If more than one value of resistance is given with different width or length values, then IRSIM will linearly interpolate the resistance based on the values. Larger than minimum length transistors are typically only used as charge keepers or in short-circuit protection at the pads, where the timing is really not relevant. If the technology does experience short channel effects (velocity saturation, etc.), and you intend to use different length

devices, you should probably generate a couple of table entries. There is no need to vary the width of the transistors since the resistance scales linearly (or very nearly so) with the width.

Note: The following part refers to files found in the source distribution of IRSIM in directory `lib/calibrate_spice3/`. Similar scripts and code can be found in the `lib/calibrate_hspice/` subdirectory.

If you have a SPICE model for your process and access to Berkeley Spice-3, you can use the **getres** script to automatically generate the resistance entries. You will first have to compile `findr.c` (use "`make -f Makefindr`"). `Getres` takes 2 arguments, a spice model, and a 'device configuration file' where you specify the sizes of the devices as well as the loading capacitance. `Getres` will setup a spice file, run spice on it, read the rawfile generated and print the table entries in a file called 'resists' as well as on stdout. The 'device configuration' file should look something like this. This one was used to generate the MOSIS 2um scalable-CMOS IRSIM prm file:

```
#
set pwith = 20.0          # with of p-device in um
set plen = 2.0           # length of p-device in um
set nwith = 10.0         # with of n-device in um
set nlen = 2.0           # length of n-device in um
set cap = 1000           # loading capacitance in fF
```

I have ratioed the p-transistor to be twice as wide to bring the threshold of the gate closer to 1/2 the swing. Although this is a crude estimate, it will work reasonably well for most processes since $n\text{-mobility} \sim 2 * p\text{-mobility}$.

You can run `getres` several times by changing the device sizes in this file.

The following files for the 2um mosis scmos process are included in the source distribution `lib/calibrate_spice3` subdirectory as examples:

```
mosis2um.spi      spice model card for the process
dev.2um          'device configuration' file used
mosis2um.prm      the prm file generated using the above
```

Obtaining parameter files for standard processes

Historically, the only open source processes were the "SCMOS" process definitions maintained by Jeff Sondeen at MOSIS/ISI. These processes are no longer supported by MOSIS. The parameter files for these processes can be found in the IRSIM source code under `lib/prm/`.

The current state of the art in open source processes are the process design kits (PDKs) produced by a collaboration of Google, Efabless, and several foundries. IRSIM support for these open-source processes can be found in the **open_pdk**s installer application, which can be obtained from either

[open_pdk on opencircuitdesign.com](http://open_pdk.on.opencircuitdesign.com)

or

[open_pdk on github.com](http://open_pdk.on.github.com)

which will install parameter files under

```
${PDK_PATH}/libs.tech/irsim/
```

Documentation for IRSIM:

Note that most of the documents below describe features of IRSIM version 9.5, and so may be outdated with respect to version 9.7. Updated versions of these documents are in progress.

The paper "Improved Models for Switch-Level Simulation" has been reproduced from a copy in the Stanford University library. Typically, these technical reports are available to the public on an FTP server. However, this particular technical report is missing from the server, so I am posting it here.

[Online \(HTML format\) Manual Pages](#)

- [IRSIM Manual](#)
- [IRSIM Analyzer Manual](#)
- [Netchange format Manual](#)

Other Online Resources

- ["IRSIM: Dynamic Power Analysis and Other Improvements"](#) by Jason Liang (September 12, 2022).

Downloadable (PDF format) Documentation

File	Date	Size	Comment
CSL-TR-88-368.pdf	November, 1988	(4.96MB)	"Improved Models for Switch-Level Simulation" (Chu)
Ackerman.pdf	March 2, 1985	(234KB)	"How to Design Simulatable CMOS Integrated Circuits" (Ackerman)
IRSIM Tutorial version 2.1.pdf	September 28, 2001	(25KB)	"IRSIM Tutorial (version 2.1)"
IRSIM_manual.pdf	Posted 10/24/06	(40KB)	IRSIM manual page
fsim.pdf	Posted 10/24/06	(30KB)	Faultsim manual page
powerEst.pdf	Posted 10/24/06	(12KB)	"Notes on Power Estimation using IRSIM"

Specifying X resources for IRSIM

The Analyzer window in **irsim** accepts X11 resource settings like all X11 applications, that can be set in a user's home **.Xdefaults** or **.Xresources** file. Some wildcard entries in a defaults file, such as **"*.background: white"**, can interfere with the analyzer display. The above example will produce a blank white-on-white display.

X11 application resources known to **irsim** are listed below, along with their default settings.

- **irsim.geometry:** *=1000x300+0+0*
- **irsim.background:** *black*
- **irsim.foreground:** *white*
- **irsim.reverseVideo:** *off*
- **irsim.highlight:** *red*
- **irsim.traceColor:** *white*
- **irsim.bannerBg:** *white*
- **irsim.bannerFg:** *black*
- **irsim.borderColor:** *black*
- **irsim.font:** *6x13*
- **irsim.borderWidth:** *2*

Troubleshooting IRSIM

Analyzer display

Symptom: The analyzer window appears with the "analyzer" command, and looks fine, but nothing shows up in the central sub-window, which is white.

Solution: If you click the mouse button in the window, and see a red line (the cursor), then the problem is that the background has been set to white and the traces are being drawn white-on-white. Usually this happens because the user has a setting like **"*.background: white"** specified in the **.Xdefaults** (sometimes **.Xresources**) file. In general, default settings like this are a bad idea,

and should be removed. But, if you have your own reasons for having such a default setting, you can specify the background specifically for IRSIM (see the section "Specifying X resources", above).

Startup

Symptom: All the signals in the window (except for the ones used as inputs) are undefined (drawn with a gray stipple pattern).

Solution: A common oversight is not issuing the commands "**h vdd**" and "**l gnd**" at the beginning of the command file. Without explicit settings, these signals are undefined and propagate the undefined status throughout the whole circuit.

Simulation

Symptom: One or more signals is undefined, but I'm pretty sure my circuit is correct.

Solution: There are various possible reasons for this, which you will probably have to check one by one.

1: First, make sure that your circuit is reset properly. IRSIM will tend to work better with flip-flops that are properly set or reset upon issue of a reset signal. Using flip-flops without a reset will work with IRSIM, but may take some extra work (see below).

2: Watch out for circuits with feedback loops onto themselves. The simplest case of this is the back-to-back inverters in a flip-flop. Using the command "**settle t**" (with *t* on the order of 50) will usually solve this problem.

3: Other circuits may have looser feedback loops that are not countered by the "**settle**" command. In the worst case, it may still be possible to "bootstrap" these circuits during a reset phase by forcing one or more nodes in the circuit low or high, as appropriate, through one or two clock cycles, then releasing the node(s) using the "**x nodename**" command. The command "**relax**" was added to version 9.7.49 to simplify this procedure, as it instantaneously attempts to resolve all undefined nodes in the circuit. In most cases, it will suffice to put the command "**relax**" at the top of the command file, thus setting initial conditions. Since this step obscures potential problems caused by random initial conditions in a circuit, it should be used with caution, and as a last resort (i.e., don't get into the habit of placing "**relax**" at the top of every command file as a precaution against undefined states).

4: Watch out for improperly defined parameter files. I had a monostable circuit fail because I used a parameter file with the wrong lambda and wrong values for source and drain capacitance, and IRSIM calculated the unit delay as under a picosecond, rounded it to zero, and the zero delay naturally prevented the monostable circuit from working regardless of the number of inverter stages used to lengthen the output pulse.

5: Finally, note that IRSIM just may not be able to handle certain circuits that are not designed with the standard CMOS methodology. I had to disable my verilog synthesizer from using the standard cell XOR and XNOR gates. In the circuit, either was easily replaced by an AOI22 gate (AND-OR-INVERT, two two-input AND gates followed by one two-input NOR gate), which can substitute for an XOR or XNOR if you have both the inverted and noninverted versions of the inputs. This is not to say that the original XOR and XNOR gates were not valid circuits, but I have not yet found a way to coerce IRSIM into handling them.

General

Have a question that is not covered here?

Try the Wiki site. Post email to me or on the Wiki site, describing the problem.

email: tim@opencircuitdesign.com

Last updated: October 2, 2022 at 4:19pm