

Netgen version 1.5 Tutorial

Table of Contents

[Disclaimer](#)
[Netgen introduction](#)
[The Netgen LVS algorithm](#)
[Generating netlist sources](#)
[Setting up for LVS](#)
[Tutorial 1: Running LVS](#)
[Tutorial 2: Interpreting LVS results](#)
[Tutorial 3: Non-matching element counts](#)
[Tutorial 4: Imbalanced hierarchy example](#)
[Tutorial 5: Handling non-matching subcell names](#)
[Tutorial 6: Imbalanced hierarchy at the top level](#)
[Tutorial 7: Non-matching component parameters](#)
[Tutorial 8: Symmetries in circuit hierarchy](#)
[Tutorial 9: Multiple devices in parallel](#)
[More setup file tricks](#)
[Tips for large netlists and messy output](#)
[Tutorial history](#)

Disclaimer

This tutorial corresponds to Netgen versions 1.4.75 and newer. The developer attempts to keep it accurate and up-to-date, but due to code fixes and new development, this cannot always be guaranteed. Feel free to contact the developer (email information at the bottom of the page) if you find anything that does not work according to the description in the tutorial text.

Netgen, a netlist manipulation and comparison tool

Netgen is a program with two purposes: One is to convert netlists between different formats, and one is to compare two netlists to determine if they are equivalent, and if not, what makes them different. The netlist conversion is a fairly trivial aspect of the netgen tool, since all netlist formats convey mostly the same information and therefore are all very similar to one another. This tutorial concerns mainly the netlist comparison. Netlist comparison is often called "**LVS**", or "Layout vs. Schematic", because its primary purpose is to check whether a VLSI layout is equivalent to the schematic from which it is derived.

A circuit designer will start with a circuit schematic of the system being designed. The circuit schematic is generally hierarchical, and starts from the fundamental devices (resistors, capacitors, transistors, etc.), and goes up to the representation of an entire chip. The program in which the circuits are drawn is normally a **schematic capture** tool as well as a schematic drawing tool. The schematic capture takes the drawn design and converts it into a netlist describing the circuit purely in forms of devices and connectivity. The netlist is typically in a well-known format such as SPICE, and is good both for simulation and for netlist comparison (although one does not necessarily use the *same* netlist for both).

The layout engineer (who may or may not be the circuit designer), will take the schematic drawing and use it to generate a layout of the circuit, while carefully considering the placement and rotation of elements, the sizing of connecting wires, and all the design rules governing the placement and sizing of geometry in the layout. Depending on the design and the tools being used, this process can be anywhere from fully automated to fully manual. But in VLSI design, especially in analog and mixed-signal (both analog and digital) designs, the layout is often generated independently of the schematic, where layout is still done mainly by hand due to the many considerations that come into play. Once the layout has been made, the layout tool (or sometimes an independent tool) will **extract** the design,

finding all the devices, determining their connectivity to other devices, and dumping a netlist containing the information. So, from the schematic and from the layout, there will be two netlists, and on some level they should represent the same circuit.

Because custom hand layout is difficult, the typical purpose of LVS is to check what errors the layout engineer has introduced in the design during the process of creating the layout. The schematic is generally assumed to be the (correct) standard, largely because the schematic is typically used for simulation, and the simulation results indicate whether the circuit (or entire chip) works or does not work as specified. This is known as *functional verification*. Even in digital designs, where the layout is automatically generated (called "**synthesis**"), LVS can be a critical help when, for example, a correction is made in the design for a metal-mask re-spin of the chip. In that case, the rewiring is usually done by hand and must be checked against a schematic to make sure it is correct. Another reason is to check that the circuit synthesis has done the job that it claims to have done. Synthesis is a very complicated process itself, and requires many setup files. It is very easy for something to go wrong; a pin placement in one file may not match the same pin placement in another file, leading the synthesis tool to wire up a circuit at the wrong point, creating an incorrect an non-functional circuit. LVS will help to pinpoint that kind of problem.

To be sure, a chip can be entirely verified through simulation, by extracting the netlist from the top-level layout and confirming that everything works through exhaustive testing of all functions. But that can take years of simulation time, depending on the size and complexity of the chip, and when the simulation fails, it is not always easy to figure out what was wrong with the layout.

Figuring out what is wrong with a mismatched layout requires some knowledge about the matching algorithm, which this tutorial aims to explain.

The Algorithm

The algorithm used by Netgen is quite elegant but can produce puzzling results if you don't understand how it works. It operates by a method called "fracturing", which is done in a series of iterative passes. On the first pass, everything in the design (both circuits, without regard to which is which) is gathered into two groups, one a list of all the devices, and the other a list of all the nets. Then, each pin of each device is given a random identifying value (such that every instance of that device in the design has the same values on the same pins). Each device type is given a unique identifying value. Each net is given an identifying number which is the *fanout* of that net (the number of device pins to which it connects).

On each pass of the device list, the list is "fractured" into multiple lists ("partitions") containing devices with the same identifying value. So after the first pass, the single device list will be broken up into individual lists for each device type. Likewise, the net list is fractured into multiple lists containing nets with the same identifying value. So after the first pass, the single node list will be broken up into individual lists for each set of nets having the same fanout.

Each device partition is given a unique identifying number, and each net partition is given a unique identifying number.

Now, the tricky part: Each device is given a new identifier that is created by combining the identifiers of each net partition and net for every net that the device connects to (this is done by summing together the results for each pin of an exclusive-or of the net partition's identifier and the net's individual identifier). Each net is given a new identifier that is created by combining the identifiers of each device partition and pin for every device pin that the net connects to (using the same method).

Repeat from the step of fracturing partitions of same-valued devices and nets. Each step greatly reduces the number of elements or nets in each partition. If the two circuits are equivalent, then the end result is that every partition contains exactly two devices or nets, one from the first circuit and one from the second circuit.

Note that in some cases, circuits may have multiple instances of indistinguishable circuits. When the circuits are truly indistinguishable from one another, they will all end up in the same partition no matter

how many iterations are made. These are called "**automorphisms**". They can be dealt with by assigning one pair of devices in the partition with a unique identifier, forcing them into their own partition, and continuing until all devices and nets are uniquely identified.

It is important to understand that while the fracturing algorithm is extremely fast and efficient, it has one significant drawback in that it does not really understand the overall structure of a circuit, and cannot produce a judgement on just how well or how badly a circuit is misconnected. There exist LVS tools that attempt to tell the designer what changes need to be made to "correct" a non-matching circuit, but such advice is all too often useless, or worse, misleading. Often it is quite difficult to figure out what has gone wrong when the error is not a simple one. The purpose of this tutorial is to show how to single out an error and ignore the extraneous information. With a few simple strategies, it is possible to work through even the most horrifically mangled design and quickly correct the problems.

Netgen can run LVS *hierarchically*. This is useful for very large circuits (like entire chips), preventing the necessity of flattening the whole circuit, producing unmanageably large output. Each subcircuit that can be clearly paired off against an equivalent subcircuit in the circuit being compared, is matched at the level of the subcircuit. If the circuits match, then the comparison moves to the next level of hierarchy, where all instances of that subcircuit are treated as individual devices with pins. If the subcircuit has no match in the other circuit, or if the comparison between subcircuits fails, then the contents of each instance of the subcircuit are pulled into the parent cell, and the subcircuit is removed. Because two netlists may define subcircuits with pins in a different order, the pin lists have to be matched up to one another. Inability to match the pins is another reason to expand the contents of the subcircuits.

Generating Netlists from Schematic and Layout

Netgen can take its source netlists from many different tools; its primary format is SPICE, and it accepts many common extensions to plain Berkeley SPICE3 syntax. It knows a handful of other netlist formats, too, but they are not so complete in their description of fundamental devices, or understand how to parse device properties, and some formats are not hierarchical.

For the purposes of this tutorial, we will work with netlists generated from a schematic using [XCircuit](#) and [Qflow](#) (where the netlists are automatically generated), and netlists generated from an extracted layout using the layout tool [Magic](#).

The tutorial will use various netlist and setup files, which can be downloaded now or as needed. Each tutorial is self-contained in its own directory. The following list has individual directories and files. The tutorials are ideally followed in numerical order.

tut1	Tutorial 1: Running LVS
tut2	Tutorial 2: Interpreting LVS results
tut3	Tutorial 3: Non-matching element counts
tut4	Tutorial 4: Imbalanced hierarchy example
tut5	Tutorial 5: Handling non-matching subcell names
tut6	Tutorial 6: Imbalanced hierarchy at the top level
tut7	Tutorial 7: Non-matching component parameters
tut8	Tutorial 8: Symmetries in circuit hierarchy
tut9	Tutorial 9: Multiple devices in parallel

In case downloading these individually is a pain, here is a tarball containing them all:

File	Revision	Size	Date
tutorial.tgz	0	(79kB)	May 19, 2016

The file "**map9v3.spice**" found in a number of the tutorials is a netlist generated by the layout tool [magic](#) from a VLSI design synthesized from verilog source using [Qflow](#) (actually the predecessor to qflow). As digital circuits go, it is relatively small, and good for demonstrating the various features of netgen (see if you can figure out what the circuit does!). The remaining files "**map9v3_testn.spice**" are hand-manipulated variations of the original netlist, used to break the netlist in ways that demonstrate how to track down and fix LVS errors. The file "**map9v3_synth.spice**" is derived independently from a netlist created by the synthesis, prior to generating the layout (placement and routing). Because the synthesis tools work on the logic gate level, not the transistor level, the netlist contains only logic gate names, and so the file "**osu035_stdcells.sp**" contains the description of the contents of each logic cell. The digital circuits are synthesized from standard cells using the [OSU standard cell set](#) that is distributed with qflow.

Setting up for LVS

The comparison algorithm only understands the topology of the netlist, the connectivity between various devices in the design. Because SPICE is a well-established netlist format, having a netlist in SPICE allows the LVS system to understand a bit more about the devices, such as being able to tell apart transistors, capacitors, and resistors. These low-level devices with pinouts established by SPICE syntax are called "**fundamental**" devices. SPICE does not go into a great deal of detail over its transistors and other devices in the statements of the netlist where they are defined, often specified by a **model** name. The name refers to a model definition, which defines the default properties of the device. The model definition is found elsewhere in a (complete, e.g., for simulation) SPICE netlist. In many systems the model is in a separate file and is expected to be appended to the netlist before simulation. This is both good and bad for netlist comparison. On the good side, if both the schematic-derived and layout-derived netlists refer to the same model names for the same devices, then it is easy to compare devices using their model names. If, however, the schematic and layout use different model names, then one must provide the netlist comparator with additional information. For this reason, the LVS uses a "setup file". In cases where the schematic-derived and layout-derived netlists use the same device model names and the same names for subcircuits, and keep the same hierarchy, everything is straightforward and no setup file is needed. Anything more complicated requires a setup file, which is described below.

Netgen has a very simple concept for the setup file. The setup file is simply a file containing commands exactly as they would appear on the netgen command line. The file is sourced like any Tcl script, so Tcl/Tk commands can be used in the setup file as well as any netgen commands.

This tutorial will first run through some simple examples that do not require a setup file, before presenting some more complicated examples that do.

Tutorial 1: Running LVS

We'll start with a simple example. Download the files in [tut1](#), including the two netlist files **bufferA.spice** and **bufferB.spice**. Like all files used in these tutorials, they are SPICE netlists. Examine both files with an editor of your choice. There are numerous differences between the formatting of the two files. **bufferA.spice** is extracted from a layout in Magic. Magic, for whatever historical reason, counts transistor devices from 1000, and adds information about the transistor area and perimeter that will increase the accuracy of a simulation. **bufferB.spice** has been created by a schematic capture tool, or possibly has been written by hand. The carriage-return characters and comment line indicate a Windows environment. It uses a non-standard (non-Berkeley SPICE, that is) keyword ".backanno". The order of the pins on the definition of subcircuit "inverter" is different from that of the first netlist. It contains device width and length information, but not area and parameter. But in spite of these syntactical differences, they describe the same circuit. A schematic drawing of the buffer circuit, showing the hierarchy down to the transistor level, is shown below.

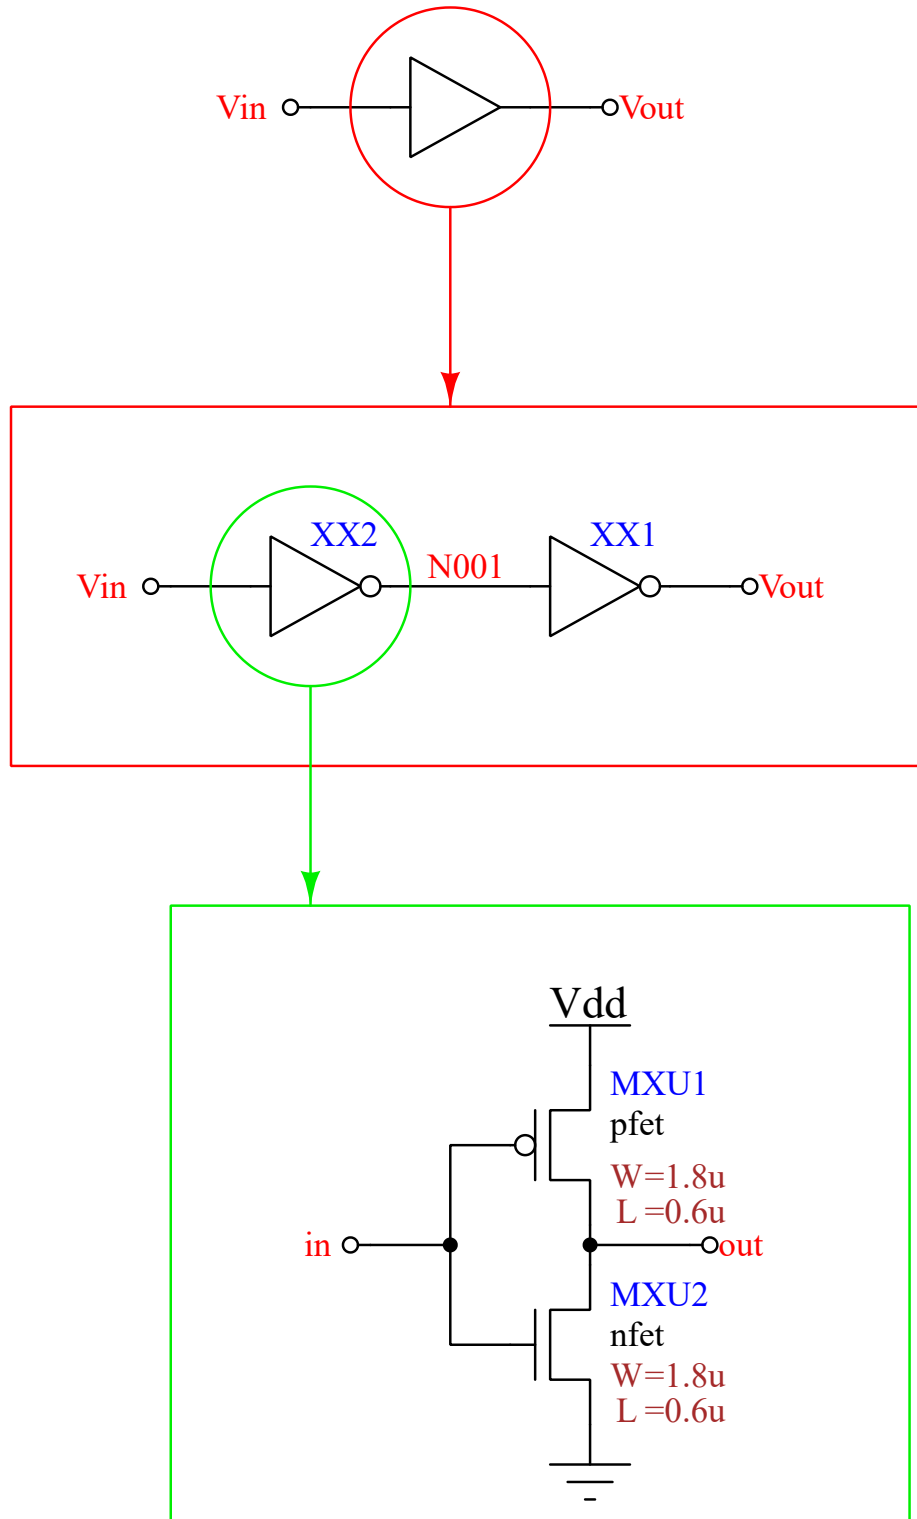


Figure 1: Schematic of the buffer circuit used in the first examples.

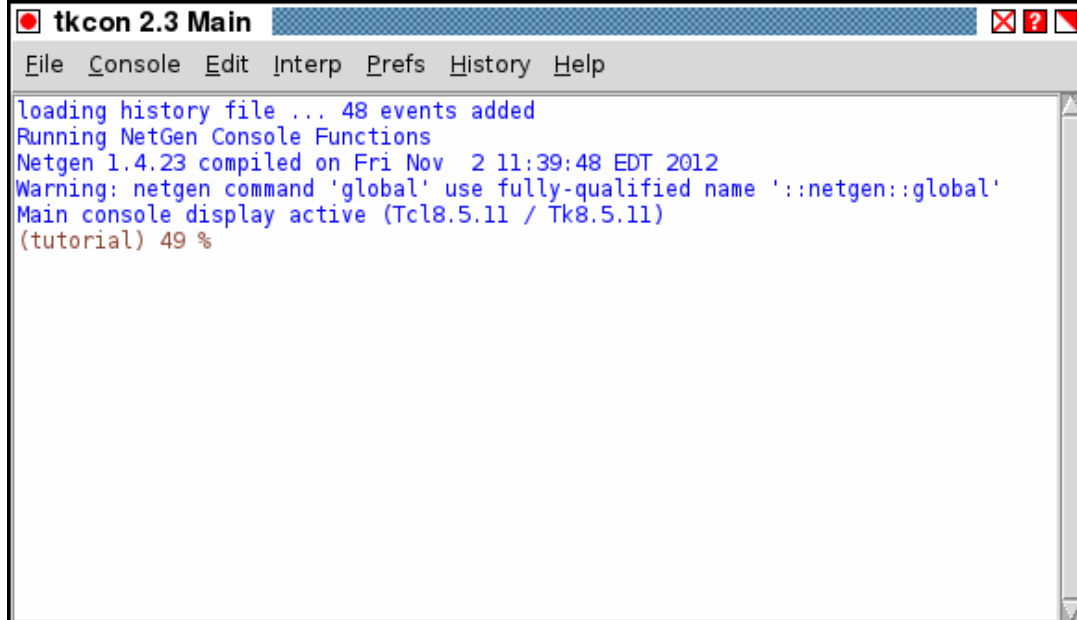
For each tutorial in this series, it will be assumed that either the tarball of tutorial files was downloaded and extracted, or the directory structure was created to follow the list above, with each tutorial self-contained in a directory called "**tut1**", "**tut2**", etc. In each directory is one or more shell script files that can be used to run netgen for each of the tutorial examples. For the first tutorial, however, we will show how netgen can be run in several different ways, both interactive and in batch.

So for the first tutorial, proceed to the directory "**tut1**". We will start with an interactive run of netgen.

Start netgen from the command line using:

```
netgen
```

You will get a console window that looks something like the following:



At the console window prompt, type the command

```
lvs bufferA.spice bufferB.spice
```

You will get a bit of output where netgen describes what it's doing as it goes along, ending with the lines

```
Result: Circuits match uniquely.  
LVS Done
```

Indicating that netgen agrees that the two netlists describe the same circuit. If you scroll back through the console output, you'll see various things, including where netgen gripes about the unknown ".backanno" command, where it shows each cell containing one nFET transistor and one pFET transistor, counts the nodes and devices, and reorders the pins to the inverter on one circuit to match the other.

The complete output will be found in an output text file dumped by netgen called **comp.out**. This is the default name used when the "**lvs**" command is run without the additional two arguments which are the name of the setup file and the name of the output log file, respectively. Edit or view this file now. In this file you will find a side-by-side comparison of the two netlists. Each one is decomposed hierarchically. First, the subcell "inverter" is presented, along with the devices it contains and the number of each. Then, there is a comparison result for *only* the inverter subcircuit. The result being a match, it is followed by a list of pins on the inverter subcircuit, matched up between the two files. Here you will see that most pins have the same name in both files, except that file "bufferA" uses the name "Gnd" for the ground node, while file "bufferB" uses the name "0" for the ground node.

After the pin lists for the inverter subcircuit, the output is repeated for the top-level circuit. The devices in each circuit are shown, along with the number of each kind. In this case, there is only one type of device, the subcircuit "inverter", and each file has two instances of each. The result of comparing the two circuits follows (which again is very short, because the circuits match exactly). Because the top-level circuit does not itself have pins, there is no pin list matchup at the end of the file for the top-level cells.

This is an example of a perfectly matched pair of netlists, and this sort of minimally informative output is the goal of any LVS run. As the number of errors between two circuits rises, the amount of output in the "comp.out" file rises, usually exponentially. Coping with the voluminous output from mismatched circuits is the main goal of the tutorial.

To end netgen, either type "quit" at the console prompt, or select "**File->Quit**" from the console pull-down menu.

Batch Mode LVS

While screenshots in these tutorials show sessions being run from the netgen console window, it is often preferred to run netgen in batch mode, and this will be the preferred way to run the tutorial

examples.

In the **tut1** tutorial directory is a script file called **runtut1.sh**. This script is all of 89 bytes long. All it does is to run netgen with the "-noconsole" option (which then uses the terminal for input, and not a console window), and pass the name of the two netlist files in the tutorial directory as arguments to the "lvs" command in netgen, just like was typed into the console in the section above. It adds the two additional arguments for setup file and output log file. No setup file exists, so this argument is quietly ignored. The output log file is specified as "**comp_tut1.out**" ("comp" being short for "comparison").

This shell script should have permissions set to executable. From the terminal's shell prompt, do:

```
runtut1.sh
```

and you will get the same result as before, except that the console output is sent directly to the terminal (and therefore can be more easily redirected, such as into another log file), and at the end of it, netgen quits and returns you to the shell prompt. This batch style of executing netgen can be used with any tutorial example not having any interactive commands.

Another way to run a batch process is from the command line, using the "-batch" option:

```
netgen -batch lvs bufferA.spice bufferB.spice
```

This form of command could be used for the shell scripts that run netgen in the tutorial examples. However, note the difference that the way the shell scripts are written call netgen directly and then feed it a series of commands. Usually this is just the "lvs" command followed by the "quit" command, but in some cases it may be useful to pass additional commands to netgen prior to, or subsequent to, calling LVS. Setup files are another way to pass commands to netgen and will be described in other tutorials.

Tutorial 2: Interpreting LVS Results

A perfectly matched netlist is all well and good, but the important thing with LVS is to understand how to interpret the output when two circuits do not match, so as to understand what to correct to make the netlists match.

Download the files from tutorial **tut2**. These include a copy of **bufferA.spice** from the first tutorial, and a new file **bufferBx.spice**. If you compare the file **bufferBx.spice**, with the file **bufferB.spice** from the first tutorial, you will see that the two files are identical except for the line ".global Vdd 0", which is missing from the file **bufferBx.spice**. Because this netlist has failed to declare nodes **0** and **Vdd** to be globals, the use of these names in the inverter subcircuit makes the nodes local to the subcircuit. The fact that there are nodes of the same names in the top level is irrelevant. These nodes define the bulk terminals of the FET devices (e.g., substrate for the nfet, n-well for the pfet).

The use of node **0** in SPICE files is generally assumed to be global. However, netgen does not make such a generalization, as it handles different formats, and other formats do not have an implicit global node **0**. If a netlist uses node **0** but does not explicitly declare it to be global, the best strategy is to set up the node **0** in netgen to be global in the setup file.

This tutorial provides two setup files, "setup_tut2b.tcl" and "setup_tut2c.tcl", whose uses will be explained presently.

For now, run netgen using the script **runtut2a.sh**. Note that this script has no corresponding setup file, and that means that the netlist file has power and ground nets defined (incorrectly) as local nodes in each subcell. The terminal output finishes with the message

```
Result: Netlists do not match.
```

The output from the script goes to **comp_tut2a.out**, which ends with the text "Netlists do not match". We know that the reason they do not match is due to the treatment of the two nodes **Vdd** and **0** as local nodes. How could we figure this out from the netgen output?

Note: There is a divergence here between behavior in Netgen 1.4 and Netgen 1.5. Netgen 1.4 tends to give up easily on mismatched pin lists and proceeds to flatten the offending circuits into each one's parent circuit. Netgen 1.5 understands that the lack of matching pins is either going to be a non-issue, if those pins happen to not be connected to anything in the parent cell, or else they are an indication of an unresolvable error, that is, one that's not going to disappear just because the circuits are flattened and absorbed into their parent circuits. Normally, flattening the circuits will just make the output harder to understand. However, in this case, mainly because the circuits are so simple, the opposite is true. So before explaining the output of tutorial 2a, take a look at the setup file **setup_tut2b.tcl**. A setup file contains any valid Tcl/Tk statements, including those that are part of the command set that netgen uses (see the "[Reference](#)" page for a list of netgen commands). This particular setup file has the commands

```
flatten class bufferBx.spice {inverter bufferBx.spice}
flatten class bufferA.spice {inverter bufferA.spice}
permute default
```

The purpose of these commands is to force the subcircuit "inverter" in each netlist to be removed and replaced by its contents (an nMOS and a pMOS transistor for each instance of the "inverter" subcircuit). Now run the script command **runtut2b.sh**. Like before, this ends in a match failure. Flattening subcircuits does not change the circuit, only its hierarchical structure. Netgen could insist that the hierarchy of both netlists be exactly the same. However, various differences between the methods used to create layout netlists vs. schematic netlists make that highly impractical at best. It is much better if the netlist comparison tool can look beyond the hierarchical structure and figure out whether or not the two netlists represent the same circuit. Netgen does this reasonably well, so subcircuits can be flattened by command at will, and should not affect the outcome of the matching. It may, however, as in this case, affect the readability of the output. Note that commands found in the setup file are executed between the time that the two netlist files are parsed, and the comparison begins. Also note that the presence of a setup file overrides certain behaviors, such as device permutation, requiring the default behavior ("permute default") to be added to the setup file.

Look at the output file **comp_tut2b.out**. There is no match between the subcircuits called "inverter" in the two files, because these subcircuits no longer exist, having been flattened by the commands in the setup file. In this output, we first get a "**Subcircuit summary:**", naming the two circuits (now at the top level, so the circuit name is the same as the filename), and a list of elements (components). Because the inverter subcircuit has been selectively flattened, and each top-level circuit had two inverters, each flattened netlist now has the contents of the two inverters, which is two pfet and two nfet devices each. This list is followed by a total count of elements and nodes (nets) in each cell. The first clue to the error is that Circuit 2 has 9 nodes, while Circuit 1 has only 5 nodes.

The main diagnostic output is what follows the circuit summary. This is split into two sections: The first is a list of "**NET mismatches**", and the second is a list of "**DEVICE mismatches**" (this nomenclature is used by Netgen-1.5. Previous versions used the more obscure terms "Illegal node partitions" and "Illegal element partitions"). All devices that could not be matched up 1:1 between the two circuits being compared will end up in one or more device mismatch lists, and all nets that could not be matched up 1:1 between the two circuits being compared will end up in one or more net mismatch lists. Each mismatch list (either device or net) will have a list of devices/nets, those on the left side belonging to Circuit 1, and those on the right side belonging to Circuit 2. It is very important to understand that there is no implied order in these lists! It should *not* be assumed that the first entry in the left-hand list is somehow equivalent to the first entry in the right-hand list. (That said, netgen could potentially use heuristic methods to attempt a sensible matchup between the two lists, but that is a project for a future release.)

Generally speaking, the device mismatch information is the least useful (which is why it is listed second), but here we will explain what it means anyway. Let's look at the first set:

Instance: inverter0/nfet1001 (drain,source) = (4,2) gate = 4 bulk = 4	Instance: inverterX1/nfetXU2 (drain,source) = (2,2) gate = 4 bulk = 1
--------------------------------------------------------------------------------	--------------------------------------------------------------------------------

The first line gives the name of the device. Because the hierarchy has been flattened, the name includes the name of the subcircuit in which that device appears. In this case, the names are the same in both circuits. The device is model "nfet", and the SPICE line calling the subcircuit named the instance "XU2". This device was inside the subcircuit called "inverter", whose instance was "X1" on the SPICE line. Each name includes both the model or subcircuit name and the instance name. This naming convention is entirely due to netgen, but is clear enough that one can easily determine what line in the SPICE file the instance comes from.

After the device name, netgen dumps a list of *fanouts* for each pin of the device. The fanout list is a count of the total number of device pins connected to that pin, *including itself* (multiple pins on the same device will be counted individually). So the minimum count is one, and if a "1" shows up in the device fanout, then you have a pin that is not connected to anything else, and that is nearly always an error. Here is clue number 2: The device mismatch list for the second circuit shows every FET device in the circuit with a fanout of "1" on the "bulk" pin, so "bulk" is unconnected on each device, which is clearly an error.

A quick explanation of the line "(drain,source) = (4,2)": FET devices have two terminals called the "source" and the "drain", where typically the source is on the side of the device closest to the supply rail (ground for nFET devices, power for pFET devices). Physically, however, in a VLSI layout, standard FET devices are completely symmetrical, and the use of the terms "source" and "drain" is largely symbolic, because they are electrically and functionally identical (some special FET devices have different properties on source and drain sides, but that is another topic). Without some knowledge about the structure of the circuit and which power supply rail is which, the circuit extractor cannot determine which end of the device may be the source, and which may be the drain. Due to the symmetry of the device, it doesn't matter which end of the device is called the "source" and which end is called the "drain". Therefore, we call these two pins "permutable". There is a command in netgen, "permute default", that is run as part of the setup script. This command tells netgen to make the two endpoint pins of all resistors and the source and drain pins of all devices called "nfet" and "pfet" permutable. With that command issued, a transistor in one circuit may be connected with the source at the top and the drain at the bottom, while the equivalent transistor in the other circuit is connected with the source at the bottom and the drain at the top, and they will still match. Later in this tutorial we will look at other uses for the "permute" command and pin permutations in general. Now, back to the output: The line "(drain,source) = (4,2)" means that on this device, "drain" and "source" are permutable pins. Netgen cannot be sure which is which, but it knows that one of them has a fanout of 4 (connects to three other device pins, therefore 4, including itself), and the other has a fanout of 2. So it writes the fanout according to the syntax shown above, to describe this ambiguity.

Now let's look at the net mismatch output, which is generally the most useful in pinpointing errors. Here, the first set starts with the output:

Net: Gnd	Net: Vdd
nfet/(drain source) = 2	pfet/(drain source) = 2
nfet/bulk = 2	
Net: Vdd	Net: 0
pfet/(drain source) = 2	nfet/(drain source) = 2
pfet/bulk = 2	

Remember what was said earlier about there being no ordering of nets within a set. Therefore it is not implied that "Gnd" in Circuit 1 matches "Vdd" in Circuit 2, only that "Gnd" and "Vdd" are two nets that do not match, and happened to end up in the same set.

The first line of each block of output is "Net:" and gives a name of the net. This is followed by a fanout, similar to the device mismatch output, except that each entry in the fanout list is a device type (the instance is not specified) and pin that connects to the named net. Here, it is very easy to see that the net "Vdd" is connected to the drain (or source) and bulk of two pFET devices in Circuit 1, while it is connected only to the drain (or source) of two pFET devices in Circuit 2. Clearly Vdd is not connected to the pFET bulk connections anywhere in Circuit 2, and that is clearly the problem. It is the last and final clue, and all clues point to disconnected bulk connections, which can be traced back to the lack of a global declaration for Vdd.

Note that in the net mismatch output, if two or more pins on a device are permutable, they are listed together, in parentheses, and separated by "|" to indicate a logical OR, as in the above output "pfet/(drain|source) = 2". This syntax means that the net is connected to 2 pfet pins, each of which may be either a drain or a source connection.

There are several ways to fix the error in the file. The best way is to make the tool that produces the netlist handle the global nodes properly. Sometimes, though, it is better just to get netgen to patch up problems with the netlists, especially if the netlist is from a 3rd party vendor, and one only has access to the netlist, not the tool that produced it. The setup file "**setup_tut2c.tcl**" is a way to get around the problem of global nodes that have not been declared global. The setup file itself declares the nodes to be global using the "**global**" command. Other than these two statements, the script "**runtut2c.sh**" runs exactly like "**runtut2a.sh**". Run the script "**runtut2c.sh**" now, and then look at the output of the file **comp_tut2c.out**". This file (and the terminal output when running the script) should show that the two netlists are equivalent. The missing "**.global**" line in the file has been taken care of by the global declarations in the setup file.

Understanding hierarhical LVS output

Now it is time to go back to the first part of the tutorial, where running the script **runtut2a.sh** produced the output file **comp_tut2a.out**. That script used no setup file, and so the result was a full hierarchical comparison.

Look at the output file **comp_tut2a.out**. In the first few lines of the file, The two inverter subcircuits are compared, and they are declared matching. Within a single subcircuit, it is impossible to know whether a node is local or global. The way netgen deals with global nodes is to treat them as local within a subcircuit, but to pass them up to the parent circuit through pins that it adds to the subcircuit. In the **comp_tut2c.out** output, following the subcircuit match, there is a list of pins for the two subcircuits and how they match up to each other. There, it can be seen that the global nodes "**Vdd**" and "**0**" of circuit "**bufferB.spice**" have been converted to ports of the subcircuit. Compare that to the output in file **comp_tut2a.out**. Instead of the global nodes being converted to pins of the inverter device, netgen declares that circuit **bufferBx.spice** has no pins equivalent to "Vdd" and "Gnd" in Circuit 1. This is the first clue, and the problem could be debugged using this information alone. However, we will continue to look at the output file.

In this output, we get two "**Subcircuit summary:**" lines, one for the inverter subcircuit, and then one for the top level. On the top level, these are no longer subcircuits, so the name given to the circuit is just the name of the file. Each subcircuit summary has a list of devices, followed by the number of instances of each device. Each top-level circuit in this tutorial example has two inverters, which is shown in the summary.

Subcircuit summary:	
inverter (2)	inverter (2)
Number of devices: 2	Number of devices: 2
Number of nets: 5 **Mismatch**	Number of nets: 9 **Mismatch**

The second clue to the error is that Circuit 2 has 9 nets, while Circuit 1 has only 5 nodes.

The main diagnostic output is what follows the circuit summary. The format of this output was covered in the previous part of the tutorial. The difference is that before, the two netlists had been flattened, and all the information about mismatches referred to the lowest level devices, the transistors. In this case, though, the low-level transistors were found to match perfectly between the two inverters, and so the circuits have been declared matching up to the level of inverters. On the top level, the devices being matched are the inverter subcircuits, not the transistors. As stated earlier, the device mismatch information is the least useful, but here we will explain what it means anyway. Let's look at the first output set in "DEVICE mismatches:"

Instance: inverter0	(no matching instance)
neg = 4	
in = 2	
out = 1	

```

pos = 4
Vdd = 4
Gnd = 4
Instance: inverter1          (no matching instance)
neg = 4
in = 1
out = 2
pos = 4
Vdd = 4
Gnd = 4

```

The first line gives the name of the device. The second column has "(no matching instance)", although the casual human observer will quickly note that the second set, where column two has entries but column one doesn't, has the obvious matching instance. Netgen only knows that it has two groups of instances that it cannot resolve, and that don't match each other in the sense of having the same fanouts on the same pins. So back to the first set.

After the device name, netgen dumps a list of *fanouts* for each pin of the device. The fanout list is a count of the total number of device pins connected to that pin, *including itself* (multiple pins on the same device will be counted individually). So the minimum count is one, and if a "1" shows up in the device fanout, then you have a pin that is not connected to anything else, and that is nearly always an error. There are no real clues in the first partition, as the "1" fanouts are for the circuit input and output. It might look strange that the inverter has sprouted six pins, but if you look at the way the netlist is written, you'll see that within the circuit the inverter has unconnected nets "pos" and "Vdd", and unconnected nets "neg" and "0". It is not uncommon to have such an arrangement on standard cells, where (in this case), "pos" and "neg" are the transistor terminals to power and ground, while "Vdd" and "0" are the well and substrate connections. They are not tied together inside the cell, so they are separate pins. "Vdd" and "0" were declared to be globals, so in the SPICE file they did not appear as pins of the device. Netgen, however, to make sure that it can properly resolve problems stemming from the use of global nets (like we're doing in this tutorial), casts every global net into a local net and brings it out as a separate pin. So we have six pins instead of four. It might also look strange that the fanouts show connections to four devices, but this is due to double-counting: The "Vdd" pin, for example, is connected to the "Vdd" input on both inverters (remember that it counts itself), and also to the "pos" pin in both inverters: That makes a fanout count of four.

The second set is more instructive, if a bit weirder looking:

```

(no matching instance)      Instance: inverterX1
                             neg = 2
                             in = 2
                             out = 1
                             pos = 2
                             proxyVdd = 1
                             proxyGnd = 1
(no matching instance)      Instance: inverterX2
                             neg = 2
                             in = 1
                             out = 2
                             pos = 2
                             proxyVdd = 1
                             proxyGnd = 1

```

The pins "proxyVdd" and "proxyGnd" may be a bit confusing. Because netgen tries to match up each level of the hierarchy pin-for-pin, it will add pins as necessary as stand-ins (i.e., a proxy) for ones that were needed for a match, but could not be found in the netlist. This is the best clue to the problem so far: Netgen had to add two pins to the inverters in Circuit2 to make them match the inverters of Circuit1. It is not quite enough information to know where to lay the blame.

Now let's look at the net mismatch output, which is generally the most useful in pinpointing errors. Here, the first partition starts with the output:

```

Net: inverter_1/in          Net: N001

```

inverter/in = 1	inverter/in = 1
inverter/out = 1	inverter/out = 1
Net: out	Net: Vdd
inverter/out = 1	inverter/pos = 2
Net: in	Net: 0
inverter/in = 1	inverter/neg = 2
(no matching net)	Net: Vout
	inverter/out = 1
(no matching net)	Net: Vin
	inverter/in = 1

It is critically important to remember, and worth repeating, that the lists are *not* aligned vertically! Do not try to match "out" on the left with "Vdd" on the right. Quick inspection of both sides shows that the nets N001, Vout, and Vin are the same on both sides. The two nets on the right side that have no match on the left are then "Vdd" and "0".

In the second set of mismatched nets, we have:

(no matching net)	Net: dummy_9
	inverter/proxyGnd = 1
(no matching net)	Net: dummy_8
	inverter/proxyVdd = 1

and in the last set of mismatched nets, we have:

Net: Vdd	(no matching net)
inverter/pos = 2	
inverter/Vdd = 2	
Net: Gnd	(no matching net)
inverter/neg = 2	
inverter/Gnd = 2	

These two sets, taken together, show exactly what is wrong with the circuit. On the left side, Circuit 1 has a net "Vdd" that is connected to both "pos" and "Vdd" pins of two inverters, and a net "Gnd" that is connected to both "neg" and "Gnd" pins of two inverters. In Circuit 2, the same nets have been split up into three parts: the "pos" pin shows up in the first set with two connections, but the "Gnd" and "Vdd" nets show up in the second and third sets with only one connection. As noted before, a fanout of one is rare and would normally be found on top-level inputs and outputs, or the occasional unused internal output (such as a standard cell flip-flop that has complementary outputs, but only one of them is needed). A fanout of one on any net that appears to be a power or ground net is almost certainly an indication of an unconnected terminal. These isolated Vdd and Gnd nets are the final bit of information needed to pin down the error in this circuit.

Getting detailed information about components and nets

Suppose that given all the clues, we still could not figure out where the error is? The **comp_tuta.out** file, although full of information about bad device and net matching, does not give a complete description of everything. For example, the devices list fanout counts on each pin, but not what net names are connected to each pin. The nets list fanout counts of devices connected to the net, but do not specify which instances of those elements are connected. These things can be queried from the netgen command line. To do that, netgen has to be run interactively (see Tutorial 1, above). As long as the netgen console window is still up after running LVS, commands may be given to the Tcl/Tk interpreter and can reference the two circuits being compared.

For this tutorial, run netgen interactively:

```
netgen
```

At the console window prompt, type the command

```
lvs bufferA.spice bufferBx.spice setup_tut2b.tcl
```

This runs LVS using the setup file from earlier in the tutorial that flattens both circuits.

In the output file, **comp.out** (the default name), we have, for example, "Instance: inverterX1/pfetXU1" in the column for Circuit 2, with "bulk = 1", which is the disassociated node. If we want detailed information about this element, then type in the console:

```
nodes inverterX1/pfetXU1 bufferBx.spice
```

Where "**bufferBx.spice**" is both the top-level cell and the filename associated with Circuit 2. The output of this node query is the following:

```
Device 'inverterX1/pfetXU1' Pins:
Pin 1 (drain) = Vout
Pin 2 (gate) = N001
Pin 3 (source) = Vdd
Pin 4 (bulk) = inverterX1/Vdd
```

This gives us all the information we need: The bulk connection on this specific element instance is "inverterX1/Vdd ". Because in this case the name "Vdd" is prefixed with an instance name, the node is *not* the same as "Vdd" by itself; it is a node local to the instance "inverterX1".

A node can also be queried to find all of the elements connected to it. If we want to know if there are any other device pins connected to that node "inverterX1/Vdd", we can find out using the command:

```
elements inverterX1/Vdd bufferBx.spice
```

which yields the output:

```
Node 'inverterX1/Vdd' in cell 'bufferBx.spice' connects to:
inverterX1/pfetXU1/bulk
```

showing that the bulk connection of "inverterX1/pfetXU1", is not connected to anything else (which was already stated by the fanout of 1 shown for this connection in the **comp.out** file).

The netlists as seen by the "nodes" and "elements" commands reflect the final state when LVS is finished. Cells that were unmatched and therefore flattened in the course of netlist comparison can also be queried. For example, use the command:

```
nodes inverter0/pfet1000 inverter bufferA.spice
```

and the result is:

```
Device 'inverter0/pfet1000' Pins:
Pin 1 (drain) = out
Pin 2 (gate) = inverter_1/in
Pin 3 (source) = Vdd (global)
Pin 4 (bulk) = Vdd (global)
```

while the command:

```
nodes inverterX1/pfetXU1 bufferBx.spice
```

yields the result:

```
Device 'inverterX1/pfetXU1' Pins:
Pin 1 (drain) = Vout
Pin 2 (gate) = N001
Pin 3 (source) = Vdd
Pin 4 (bulk) = inverterX1/Vdd
```

This is yet another way that shows that the bulk connection to the transistor is a local node, where it should have been a global node or a port to the cell.

Queries can also be made on the original cell hierarchy, even though it has been flattened. If you do this, though, be aware that the original cell is no longer part of the network participating in LVS. So this command:

```
nodes pfet1000 {inverter bufferA.spice}
```

specifies the device "pfet1000" in cell "inverter" of file "bufferA.spice". Instances of the cell "inverter" were flattened for LVS, but the cell still exists in memory, so the above command yields the result:

```
Device 'pfet1000' Pins:  
Pin 1 (drain) = out (port of inverter)  
Pin 2 (gate) = in (port of inverter)  
Pin 3 (source) = pos (port of inverter)  
Pin 4 (bulk) = Vdd (global)
```

Tutorial 3: Non-Matching Element Counts

Let's look at the output for the case of non-matching element counts. Get the files of **tut3**, which are similar to the netlists used for the previous tutorial. However, "**bufferB.spice**" has been modified. Before the ".end" statement, the following line was added:

```
MXU3 Vdd Vdd Vdd Vdd pfet W=1.8u L=0.6u
```

This is a typical entry for a "dummy" transistor, one which is added to a layout to provide better matching of the characteristics of different transistors in an array, or in a differential pair. It is not uncommon to forget these in a schematic drawing, so it is useful to know what they look like when they cause a netlist to be unmatched. Run the script "runtut3.sh" to make the comparison of the two netlists.

Assuming you did not know beforehand what change was made to the second netlist. In the **comp_tut3.out** output file, there is a clear clue to the error. In the summary of the comparison of the "inverter" circuits:

Circuit 1: inverter	Circuit 2: inverter
pfet (1)	pfet (2) **Mismatch**
nfet (1)	nfet (1)
Number of devices: 2 **Mismatch**	Number of devices: 3 **Mismatch**
Number of nets: 6	Number of nets: 6

Here we have a very clear picture of what's going on: The summary shows that both circuits have cells named "inverter", but the "inverter" of Circuit 2 has the additional "pfet" component that does not exist in Circuit 1. Once again, this error plays out in both the net mismatch list and the device mismatch list. The device mismatch list shows the extra pfet, with "(no matching instance)" on the other side for Circuit 1, and the name of the instance ("pfetXU3") is shown, which leads us to look for a line "XXU3 ... pfet ..." in the netlist file. And from the net mismatch list, we can see that the net **Vdd** has connections to pfet drain (source), gate, and bulk, consistent with a dummy transistor.

In general, the best way to work through LVS problems is to start with the device count mismatches in each circuit. Once both circuits have a matching number of devices on all levels of the hierarchy, solving any remaining connectivity problems will become much easier.

It is left as an exercise for the reader to fix the second netlist to match the first, although normally this would be done in a schematic drawing environment.

Tutorial 4: A Imbalanced Hierarchy Example

It was mentioned in the first example that the reason for selectively flattening subcircuits in a netlist is to correctly identify matching in circuits with an "imbalanced hierarchy", that is, in which the hierarchical structure of one circuit is different from another, even though they describe the same netlist.

Download the files in **tut4** for this example. Files **hier1.spice** and **hier2.spice** show what is meant by an "imbalanced hierarchy". The circuit is a simple buffer, like the previous examples. If you look at **hier1.spice**, however, you will see that it has an additional layer of hierarchy by defining a subcircuit

called "buffer" and then instantiating it at the top level. File **hier2.spice** also defines a subcircuit "buffer" and instantiates it at the top level, but it does not define an additional subcircuit called "inverter", but instead implements the two inverters implicitly as a circuit of four transistors.

Run the example by running the script "**runtut4.sh**". You can see that netgen declares the two circuits matched. If you look at the output in "**comp_tut4.out**", you will see that there is no comparison at all of the subcell "invert". At the top of the file is the line:

Flattening unmatched subcell invert in circuit buffer (2 instances)

As netgen reads in two netlist files, it tries to match up subcircuits by name, on the assumption that a schematic and layout will keep the same naming conventions of an entire circuit hierarchy (apart from the top-level cell). Usually this is a good assumption. What to do when it is not a good assumption will be covered in the next tutorial. In this case, though, netgen finds a subcircuit named "buffer" in both netlists, and assumes they should be matched. However, it finds a subcircuit called "invert" in one netlist but not in the other. Having nothing to match it to, it selectively flattens the invert subcircuit, quietly, then continues with the comparison of the buffer circuits.

Tutorial 5: Handling non-matching subcell names

Sometimes cells that are functionally the same are named differently between schematic and layout. Often this is simply a matter of one being lower case and the other upper case; in SPICE, which is case-insensitive, these would be considered the same name. Netgen will ignore problems of case-sensitivity if both cells being compared have come from SPICE netlists. But sometimes, a cell or a fundamental device may just be named differently between the schematic and layout.

Download the files in **tut5**. Consider the two netlists **map9v3.spice** and **map9v3_test2.spice**, which are considerably larger than previous examples in this tutorial set. The cell **map9v3.spice** will be used for the next three examples. It is a standard-cell layout, originally generated from verilog source code. For this example, both netlists were generated from the layout, and the second one has been manipulated by hand for the purpose of the tutorial.

First, take a look at these two files in an editor of your choice. They are nearly identical, but the second netlist has replaced all the cell names with different names, and an error has been introduced into the file.

Note: The original version of this tutorial used lowercase vs. uppercase names. Because the SPICE netlist format is not case-sensitive, netgen code was reworked in version 1.4.55 (February 2014) to do case-insensitive name matching when both input files are in SPICE format. This version and newer versions of netgen will match a cell named, for instance, "NAND2" in one netlist to a cell named "nand2" in the other netlist.

Run the first part of the tutorial example by running script "**runtut5a.sh**". Note that there is no setup file corresponding to this script. The result of this run is "**Netlists do not match**". Now, look at the file **comp_tut5a.out**. At the top of the file you will find a long list where every cell has been flattened, in each netlist. What happened was that netgen did not consider any of the subcircuits in one netlist to match any of the subcircuits in the other netlist, because it could not find any names that matched. Not willing to match up the subcircuits by guessing, it flattened them all before the comparison, and made the netlist comparison at the transistor level. You will see at the end of the file a summary of the contents of the two top-level netlists, showing the netlists not as a collection of a number of standard cells, but as a collection of a large number (668) of transistors. Because the error introduced into the second file was a trivial change to an internal net, and because the hierarchy is only two-deep, the output is fairly easy to interpret and the error easy to trace. But imagine trying to trace a complicated error through a circuit of tens of thousands of transistors!

By the way, it should be clear from the **comp_tut5a.out** output that there is one network in file **map9v3.spice** called "**OAI22X1_3/Y**", and it quite clearly matches two nets in **map9v3_test2.spice**, and so it is pretty simple to conclude that the second netlist has (purposefully!) chopped one of its nets in

two, and those two nets need to be connected back together (because this is a hand-edited netlist, the solution to fix the netlist is to rename the net "**BADNODE**" to "**oai22xA_3/Y**").

So netgen isn't matching the cell names. What do we do about it? This is where the *setup file* becomes very useful. Take a look at the setup file **setup_tut5b.tcl**. You will see that it is simply a set of commands "equate classes" and each line specifies a cell in the naming convention of the first netlist, and a cell in the naming convention of the second netlist. This tells netgen that the two cells indicated in the command should be considered *potentially* equal. This will cause the cells to be compared against each other, as opposed to being flattened into the parent cell.

Run netgen with the command "**runtut5b.sh**", which uses the setup file **setup_tut5b.tcl**. Look at the resulting file **comp_tut5b.out**. The first thing to notice is that a comparison has been done separately for each of the standard cells, and each standard cell passes LVS. In particular, search for the word "Flattening". You won't find it. This indicates two things: that netgen has found a match for each subcell in each circuit, and that the error (assuming you didn't fix it) is not inside a subcell (which would cause the subcell not to match, and therefore force it to be flattened). At the end of the file, the contents of each netlist is a collection of standard cells, not transistors. The error is, of course, the same, with the same solution, but it is clear from the output that the error is on the top level of the hierarchy, and not inside a subcell, a fact that helps in finding and correcting the error.

Here's a good exercise to try: (1) correct the error in the netlist and get netgen to generate a "Graphs match uniquely" message, and (2) introduce an error (by hand) into a subcell (not the top level) and see how this affects the output. For example, take any random node name in one of the subcells in **map9v3_test2.spice** and change it to "BADNODE2". You will probably notice immediately that the error output in file **comp_tut5b.out** is much more extensive than before. That's because where a subcircuit is incorrect, that failure is replicated in every instance of the subcircuit. When netgen can't successfully compare the subcircuit in the two netlists, it flattens every instance, replicating the error in the top level cell. When tracing such an error, you will need to find the first occurrence of " NET mismatches" in the output file. Noting that this output is produced for one of the subcells, it indicates that the error in the subcell should be attended to (and corrected) first. That will avoid the hassle of having to deal with the copious output for the top-level cell comparison.

If you are looking first at the output for the top-level cells, you would be well advised to notice that the low-level devices "**nfet**" and "**pfet**" appear in the summary for both circuits. Because both circuits are made from standard cells, the presence of these low level devices should be an alert that some error occurred in a subcell and the subcell has been flattened.

Tutorial 6: Imbalanced hierarchy at the top level

Here is a not-so-uncommon issue, and a more realistic example than hand-edited netlists. Download the files from **tut6**, containing the netlist file "**map9v3_synth.spice**" and the library file "**osu035_stdcells.sp**". In the course of synthesizing a digital circuit from verilog (hardware description language), the synthesis process first generates a netlist, then generates a layout. We have previously looked at a netlist that was created by taking the synthesized layout and extracting a netlist from it. Here, we have a circuit that is converted directly from the netlist created by the synthesis tool. Run the script command **runtut6a.sh**. Netgen will produce an output file **comp_tut6a.out** file, comparing cells from bottom up. But on the console or terminal, the last thing it prints out is various warnings such as:

```
Contents of circuit 2: Circuit: 'map9v3_synth.spice'  
Circuit map9v3_synth.spice contains 0 device instances.  
Circuit contains 0 nets.
```

```
Circuit map9v3_synth.spice contains no devices.  
Result: Verify: cell map9v3_synth.spice has no elements and/or nodes. Not checked.
```

This may seem a bit odd, because the cell "**map9v3_synth.spice**" seems to have a lot of stuff in it. It's definitely not an empty netlist. But if you look closely at the file, you may notice that while the circuit contains a *subcircuit* called "map9v3", it does not actually instantiate this circuit anywhere. That is, there is no line in the file beginning with "X" and ending with "map9v3". The file is very much like the library

file that it includes ("osu035_stdcells.sp"), in that it defines subcircuits but does not contain (that is, instantiate) any actual circuit itself.

But what we really want to do here is to compare the top-level circuit in the file "**map9v3.spice**" against the subcircuit defined in the file "**map9v3_synth.spice**". One way to do this would be to strip the lines ".subckt ..." and ".ends" out of the second netlist file. But there is an easier way to make the comparison without changing the contents of the file. Run netgen again with the script "**runtut6b.sh**". This time you will get a proper comparison. The difference is in the scripts, so look at the two script files **runtut6a.sh** and **runtut6b.sh**. The difference is that the first one runs the command:

```
lvs map9v3.spice map9v3_synth.spice
```

while the second one runs:

```
lvs map9v3.spice "map9v3_synth.spice map9v3"
```

The second form tells netgen to compare against the subcircuit called **map9v3** inside the file **map9v3_synth.spice**. It does not matter if the file is a library or a regular netlist, as long as the specified cell can be found somewhere in it.

You will find out from the output of **runtut6b.sh** that the two netlists don't match. Can you figure out why?

First hint: Look at the first two nets listed in the "Net mismatches".

Second hint: Synthesized logic circuits will be placed in rows, and somebody needs to tell the router to connect up power and ground between all the rows. . .

But the main point of this exercise is to show that the use of a (Tcl) list of two names "*filename cellname*" tells netgen to make the comparison against a specific subcircuit named *cellname* that is in the file *filename*. The subcircuit could, in fact, be a circuit deeper in the hierarchy. You could, for example, start up netgen and issue the command: "lvs "map9v3.spice BUFX2" "map9v3_synth.spice BUFX2"" and you would get a comparison only between the two subcells named "BUFX2".

NOTE: In the Tcl "tkcon" console window, the "*filename cellname*" list can also be specified in braces instead of quotes, as this is an alternative way to specify lists in Tcl/Tk: {*filename cellname*}. However, for the batch file execution, the braces have a different meaning to the shell, and so quotes need to be used.

One interesting question to ask is what happens if the standard cell definitions are missing from the file? This is true of netlist file "**map9v3_nolib.spice**", which is exactly like "**map9v3_synth.spice**" except that it is missing the line at the top that includes all of the standard cell definitions from file "**osu035_stdcells.sp**". Run this example using "**runtut6c.sh**" and look at the result "**comp_tut6c.out**". Each standard cell complains about disconnected nodes and lack of devices in the cell. However, it goes ahead and does the top-level matching as before, and comes to the same result. In most circumstances this should be considered a practice to avoid, because there is no guarantee that the subcells match.

Just like we used the setup file back in tutorial 2 to correct problems with the netlist (there, with the naming of global nodes), the setup file can be used in this case to provide the missing library. Take a look at the setup file "**setup_tut6d.tcl**". It is somewhat complicated to provide the missing library file. The circuits have already been read in by the time the setup file is sourced, so the instances of all the standard cells will exist as empty placeholders. Netgen must be told to use the library file to replace the placeholders with the complete subcircuits. So the first two lines,

```
set flist [canonical map9v3_nolib.spice]
set filenum [lindex $flist 1]
```

get the file name and number of the netlist of interest, and pull out the file number. This file number is passed to the command that reads the library, so that the standard cells are attached to the specified netlist:

```
readnet spice osu035_stdcells.sp $filenum
```

The final two statements work around a completely different problem. If you look at the output "**comp_tut6b.out**", you'll find that for each standard cell comparison, the circuits match, but this is followed by a bunch of output like this:

```
Property errors were found.  
Netlists match uniquely.  
There were property errors.  
nfet1002 vs. nfet2:  
ps circuit1: 6e-06 circuit2: 0  
as circuit1: 2e-12 circuit2: 0  
pd circuit1: 1.04e-05 circuit2: 0  
ad circuit1: 4.4e-12 circuit2: 0
```

and so forth. These properties are standard parameters for transistors in SPICE netlists, where "ad" means "area of drain", "as" means "area of source", "pd" means "perimeter of drain", and "ps" means "perimeter of source". Ideally, these would be matched between the two circuits. However, layout extraction of a netlist cannot always make a definitive estimate of these values, as it has to make sure that it does not double-count, and drains and sources may overlap or otherwise complicate the extraction process. In some cases the layout extraction will combine the lump sum of the values onto one transistor node and leave the others zero, which is the case here. So it is often desirable to have netgen ignore these values instead of treating them as properties to compare, and fail if not matching. The setup file statements

```
property {map9v3.spice nfet} remove as ad ps pd  
property {map9v3.spice pfet} remove as ad ps pd
```

Will remove these values from the netlist. It does not need to be done on both netlists. If netgen sees a property in one file, but that property is missing from the other file, it will ignore both. Certain critical properties, such as length and width of transistors, will not be ignored regardless, and will generate an error message if missing. Secondary properties such as these, however, will simply be ignored unless they appear in both netlists.

Finally, it is possible to treat any subcircuit as a "black box" circuit. In essence, that is what happened with output "**comp_tut6c.out**". No subcircuit contents were found, so the subcircuits were treated as "black boxes", but with lots of error messages. The error messages can be suppressed by explicitly declaring each subcircuit to be a black box circuit. This is done in file "**setup_tut6e.tcl**" using statements such as, e.g.,

```
model {BUFX2 map9v3_nolib.spice} blackbox
```

Note that this must be done on both netlists so that all subcircuits will be properly treated as black boxes. This example can be run with the command "**runtut6e.sh**", and the output appears in "**comp_tut6e.out**".

Tutorial 7: Non-matching component parameters

The previous tutorial had some warnings involving properties (source, drain area and perimeter) that were in one file but not the other. These needed to be ignored. However, every low-level device has specific properties that define it, such as length and width for transistors, resistors, and capacitors. Many analog circuits will fail if the component values are incorrect, so it is critical to know that the component values match between the two netlists.

Download the files of **tut7**. Run the script "**runtut7.sh**". Here is what you should see at the end of the terminal output:

```
Result: Circuits match uniquely.  
The following cells had property errors: BUFX2  
LVS Done
```

The way to interpret this output is that the *structure* of the two netlists was the same, but that some element or elements within the netlist (in this case, inside the cell **BUFX2** as indicated) had properties that were different between the two netlists. A *property* is a value passed to a device, such as width or length for a transistor, a component value for a resistor or capacitor, and so forth. Netgen will parse most standard parameters like transistor width and length, and parameters whose names match between the two netlists. Netgen will ignore parameters that appear in one netlist but not in the other, assuming that parameters which are missing are not of interest to the netlist comparison.

Look at the file **comp_tut7.out** resulting from this run of netgen. The notice of property errors on the console window is a cue to go looking for the text "Property errors" in the **comp_tut7.out** file. The list of cells containing property errors is printed at the end of the output, since property mismatches don't prevent a successful network topology comparison. Otherwise, the end result would say that the netlists match, and it would be difficult to figure out if there were property mismatch messages buried in the output somewhere. If you look at the output for the comparison of subcircuit BUFX2, you will see the full explanation of the error:

```
Circuits match uniquely.  
Property errors were found.  
Netlists match uniquely.  
There were property errors.  
nfet1003 vs. nfet4:  
W circuit1: 4e-06 circuit2: 2e-06 (delta=66.7%, cutoff= 0%)
```

This pinpoints the error as being a transistor "nfet1003" in cell "BUFX2" in circuit "map9v3.spice" having a width of 4um, while the (structurally) matching transistor "nfet4" in cell "bufX2" of circuit "map9v3_test.spice" has a width of 2um. The remaining text says that the difference is 66.7%, and the netlister has a cutoff criterion of 0%, meaning that widths of the devices in each cell must be exactly the same. The difference value may be confusing, as one width is double the other, so you might wonder where the difference of 66.7 percent comes from. To make the comparison symmetric, so that the same answer is given when circuit 1 is compared to circuit 2 as for when circuit 2 is compared to circuit 1, and to prevent unbounded results when one of the values is zero, the equation used for property comparison is

$$2 * \text{abs}(\text{property}(1) - \text{property}(2)) / (\text{property}(1) + \text{property}(2))$$

This may not be a particularly meaningful metric, but by being symmetric, the cutoff criterion is the same no matter which circuit is first or second in the comparison.

Once found, the property errors can be traced back to the netlists; device "nfet1003" is, in the SPICE netlist, transistor record "M1003" with device model name "nfet", at line 10 of file "map9v3.spice":

```
M1003 Y a_10_30# gnd Gnd nfet w=4u l=0.4u  
+ ad=4p pd=10u as=0p ps=0u
```



while device "nfet4" is transistor record "M4" with device model "nfet" at line 7 of file "map9v3_test.spice":

```
M4 Y a_10_30# gnd Gnd nfet w=2u l=0.4u
```

Note that the file "map9v3.spice" has additional information for each transistor, including area and perimeter values for transistor source and drain. Because these appear in one file but not the other, netgen ignores them and does not flag this as an error.

Tutorial 8: Symmetries in circuit hierarchy

Automorphisms were described above in "[The Algorithm](#)" section as portions of a circuit which have the same connectivity and are therefore indistinguishable from one another. It is pretty easy to construct a circuit that has symmetries. For example, download the files from **tut8** and inspect the two files **morph1.spice** and **morph2.spice**. The contents can be schematically shown as in the figure below:



As you can see, the two circuits are equivalent, but the inputs are connected to different inverters in the same cell. As far as the netlist is concerned, the three inverters are identical, and indistinguishable. This makes the netlist of the two subcells "triinvert" have multiple symmetries. If you run LVS on these two files, the terminal output will report:

```
Circuits match with 4 symmetries.
Netlists match uniquely.
Circuits match correctly.
```

However, when it comes time to compare the two top-level cells, the symmetries are fixed; this happened at the time of the output line "Netlists match uniquely". Symmetries imply non-uniqueness. But the two subcircuits in the two netlists have pins, and those pins have unique names that break the symmetry. Above the level of that cell in the hierarchy, the pins "in1" and "out1" are treated as specific pins of the subcircuit "triinvert", and cannot be viewed as permutable with "in2" and "out2", or with "in3" and "out3". (Side note: The "permute" command would not be useful in this context, as it only works for individual pins, not groups of pins. Making all the "in" pins permute and all the "out" pins permute would be completely wrong, as netgen would no longer be able to see the difference between a signal that entered on "in1" and exited on "out1" in one netlist but entered on "in1" and exited on "out2" in the other.)

So these two circuits use the three inverter resources in the subcircuit "triinvert" differently, and so the netlist fails. Once this has been determined as the failure mechanism, however, it is relatively easy to fix, by flattening all instances of "triinvert" so that the pins which break the symmetries are removed. The setup file "**setup_tut8b.tcl**" does this flattening, so that if you run the script "**run_tut8b.sh**", the symmetry is kept up to the top level and the two circuits are declared matching.

It would be preferable if netgen could detect such group symmetries and selectively flatten accordingly. Unfortunately it does not do so at this time, but it is a topic for future development.

Tutorial 9: Multiple devices in parallel

Download the set of files in **tut9**. If you look at the two files to compare, **x.spice** and **y.spice**, you will see that **x.spice** has two device lines, one of which uses the SPICE parameter "M = 2" to declare two parallel copies of the device, while the file **y.spice** has two parallel devices explicitly declared.

Run the script **run_tut9.sh** and look at the output file, **comp_tut9.out**. Note that in the subcircuit summary, both sides claim to have two PMOS devices. What netgen does is to expand out any use of "M =" in a netlist into individual devices.

At this time, netgen is unable to sort out differences in number of devices that may be a tradeoff with a property such as width and/or length. That is a topic for future development.

More setup file tricks

There are occasions when a correct netlist will not be able to pass LVS due to unexpected mismatches in the hierarchy. For example, a resistor network in two cells may contain a subcircuit with four resistors (let's call it "res_array") and an additional lone resistor not in a subcircuit. Both circuits have the same

name ("res_array") for the subcircuit containing the four resistors, so those two subcircuits are matched and pass LVS. However, each netlist makes different connections to those resistors. To the LVS tool, each circuit has one instance of "res_array" and one resistor device, so they appear to be matched; but these have different network connectivity in each cell, so the LVS fails. Depending on the circumstance, the end-user may want this situation to fail LVS, such as if the resistors in "res_array" are in a special layout with best-matching practices, and using the lone resistor instead of one in the array would be considered an error.

In the case that it is known not to be an error, the LVS can be made to pass by specifically forcing netgen to flatten the resistor array in both circuits, in setup.tcl:

```
flatten parent_cell {res_array file1}  
flatten parent_cell {res_array file2}
```

Note that this use requires netgen version 1.5.53 or newer. Previous versions would only allow "flatten {res_array file1}", which would flatten *all* instances of res_array in the whole file, not just within the cell parent_cell . That works, but at the expense of losing a level of hierarchy in cells where it was not necessary to do so.

In general, properties are compared where both netlists define them for a cell, instance, or (low-level) device. If only one of the netlists defines a property, it is assumed to be meaningful only to that netlist (e.g., relevant to layout or schematic, but not both), and no warning is issued about missing properties in the other netlist. Sometimes, however, properties exist in both netlists but are often in disagreement not due to an error in the netlists, but due to the way different tools generate or use those values. A typical example is the source and drain area and perimeter values for MOSFET devices. In SPICE files, these are given the (fixed) names "AD", "PD", "AS", and "PS". In cases where these are present in both netlists, are often conflicting, and are generating large amounts of unnecessary output, they can be suppressed by deleting them from one of the netlists:

```
property {file1 nmos} remove ad as pd ps  
property {file1 pmos} remove ad as pd ps
```

where "file1" is the name of one of the netlist files. It is not necessary to remove the parameters in both netlists, although it is necessary to have one statement per device that defines those properties.

Tips for large netlists and messy output

It is generally best to avoid flattening netlists unless absolutely necessary. Use hierarchical netlists whenever possible, and keep the hierarchies the same between the two netlists. It will be much easier to track errors when each cell has a handful of subcircuits with pins instead of possibly thousands of devices. If a circuit has to be flattened, use commands in the setup file to flatten only those circuits for which it is needed (see above). Don't flatten any cells until it becomes apparent that LVS will fail without it.

Tutorial History

November 24, 2007

Placeholder created for netgen tutorial

October 8, 2012

Finally getting around to writing the tutorial. . .

November 1, 2012

Moving on to describing the examples.

May 19, 2016

Tutorials revised, improved, and expanded.

